

# Inline Caching in JavaScriptCore

Filip Pizlo  
Apple Inc.



[webkit.org](http://webkit.org)

<https://svn.webkit.org/repository/webkit/trunk>

JavaScriptCore.framework



Safari

# Agenda

- JavaScript execution strategy
- Focus on *property access* inline caches
  - Simple inline caching
  - Inline Caching Optimizations
  - Lots of perf data


# JavaScript execution strategy

# Four Tiers




↔ *latency*                                    *throughput* →

# Profiling



# Speculation and OSR





Parser

Parser

*AST*





Parser


*AST*


Bytecompiler


*unlinked  
bytecode*















































# Inline Caching

# Inline Caching

- Performance
- Why?
- Simple inline caching
  - in interpreter
  - in JIT
  - prototype
- Advanced topics
  - Polymorphic Access JIT
  - Inlining inline caches
  - Type inference

# Performance

Benchmark	Summary	Number of Samples in my experiments
JetStream 1.1	ES5 benchmarks (real and synthetic)	9
ARES-6 1.0.1	ES6 benchmarks (real and synthetic)	24
Speedometer 2.0	DOM framework benchmarks (real)	30

Source: [browserbench.org](http://browserbench.org)



JetStream

ARES-6

Speedometer 2

1×

1.5×

2×

2.5×

3×

3.5×

4×

4.5×

5×

5.5×

6×

6.5×

**Baseline = no property inline caching**


JetStream

ARES-6

Speedometer 2


1x 1.5x 2x 2.5x 3x 3.5x 4x 4.5x 5x 5.5x 6x 6.5x

**Baseline = no property inline caching**





{x: 1, y: 2}




{x: -5, y: 7}




{x: 42, y: 3}


```
var x = 0.x;
```




`O.X = X;`



# Hashtable




# Hashtable




- Pointer chasing is slow

# Hashtable




- Pointer chasing is slow
- Hash codes take time to compute

# Hashtable



- Pointer chasing is slow
- Hash codes take time to compute
- Lots of instructions, hard to inline

# Hashtable




- Pointer chasing is slow
- Hash codes take time to compute

Lots of instructions, hard to inline




{x: 1, y: 2}




{x: -5, y: 7}




{x: 42, y: 3}



## *structure*




***structure***




▶ hashtable

## ***structure***




- ▶ hashtable
- ▶ maps name to offset


## ***structure***




# JSC Object Model




# JSC Object Model




# JSC Object Model



# JSC Object Model




# JSC Object Model



# JSC Object Model

*statically configurable*



# JSC Object Model

*statically configurable*



*dynamically configurable*

# Fast JSObject

```
var o = {f: 5, g: 6};
```

structure ID: 42	Indexing	type	flags	cell state	null	0xffff000000000005	0xffff000000000006
---------------------	----------	------	-------	------------	------	--------------------	--------------------

# Fast JSObject

```
var o = {f: 5, g: 6};
```



Structure  
Table

# Fast JSObject

```
var o = {f: 5, g: 6};
```



Structure  
Table

Structure

# Fast JSObject

```
var o = {f: 5, g: 6};
```




Structure  
Table

Structure

Property Table


# Fast JSObject

```
var o = {f: 5, g: 6};
```




# Fast JSObject


```
var o = {f: 5, g: 6};
```



```
var o = {f: 5, g: 6};
```




```
var o = {f: 5, g: 6};
```




```
var o = {f: 5, g: 6};
```


```
var v = o.f;
```



```
var o = {f: 5, g: 6};
```



```
var v = o.f;
```




```
if (o->structureID == 42)  
    v = o->inlineStorage[0]  
else  
    v = slowGet(o, "f")
```

# “Inline Cache”

```
var o = {f: 5, g: 6};
```

```
var v = o.f;
```



```
if (o->structureID == 42)
    v = o->inlineStorage[0]
else
    v = slowGet(o, "f")
```

# Interpreter Inline Cache

get\_by\_id <result>, <base>, <propertyName>

# Interpreter Inline Cache

```
get_by_id <result>, <base>, <propertyName>,  
<cachedStructureID>, <cachedOffset>
```

# Interpreter Inline Cache


```
get_by_id loc42, loc43, "g",
          0, 0
```

# Interpreter Inline Cache

```
get_by_id loc42, loc43, "g",
0, 0
```



# Interpreter Inline Cache



# JIT Inline Cache

```
0x46f8c30b9b0: mov 0x30(%rbp), %rax  
0x46f8c30b9b4: test %rax, %r15  
0x46f8c30b9b7: jnz 0x46f8c30ba2c  
0x46f8c30b9bd: jmp 0x46f8c30ba2c  
0x46f8c30b9c2: o16 nop %cs:0x200(%rax,%rax)  
0x46f8c30b9d1: nop (%rax)  
0x46f8c30b9d4: mov %rax, -0x38(%rbp)
```

# JIT Inline Cache

```
0x46f8c30b9b0: mov 0x30(%rbp), %rax
0x46f8c30b9b4: test %rax, %r15
0x46f8c30b9b7: jnz 0x46f8c30ba2c
0x46f8c30b9bd: jmp 0x46f8c30ba2c
0x46f8c30b9c2: o16 nop %cs:0x200(%rax,%rax)
0x46f8c30b9d1: nop (%rax)
0x46f8c30b9d4: mov %rax, -0x38(%rbp)
```


# JIT Inline Cache

```
0x46f8c30b9b0: mov 0x30(%rbp), %rax
0x46f8c30b9b4: test %rax, %r15
0x46f8c30b9b7: jnz 0x46f8c30ba2c
0x46f8c30b9bd: jmp 0x46f8c30ba2c
0x46f8c30b9c2: o16 nop %cs:0x200(%rax,%rax)
0x46f8c30b9d1: nop (%rax)
0x46f8c30b9d4: mov %rax, -0x38(%rbp)
```


# JIT Inline Cache

```
0x46f8c30b9b0: mov 0x30(%rbp), %rax
0x46f8c30b9b4: test %rax, %r15
0x46f8c30b9b7: jnz 0x46f8c30ba2c
0x46f8c30b9bd: cmp $0x125, (%rax)
0x46f8c30b9c3: jnz 0x46f8c30ba2c
0x46f8c30b9c9: mov 0x18(%rax), %rax
0x46f8c30b9cd: nop 0x200(%rax)
0x46f8c30b9d4: mov %rax, -0x38(%rbp)
```


# “Self” IC speed-up




# “Self” IC speed-up



# “Self” IC speed-up



# “Self” IC speed-up



# Simple Inline Caching

- Interpreter edits instruction operands.
- JIT reserves a *nop sled* and patches it.

# Prototypes

# Prototypes

```
function Foo(f)
{
    this.f = f;
}

Foo.prototype.getF = function()
{
    return this.f;
}
```

# Prototypes

```
function Foo(f)
{
    this.f = f;
}
```

```
Foo.prototype.getF = function()
{
    return this.f;
}
```



getF:  $\lambda$


**Foo.prototype**

# Prototypes

```
function Foo(f)
{
    this.f = f;
}

Foo.prototype.getF = function()
{
    return this.f;
}

var o = new Foo(42);
```




# Prototypes

```
function Foo(f)
{
    this.f = f;
}
```

```
Foo.prototype.getF = function()
{
    return this.f;
}
```

```
var o = new Foo(42);
```

```
var tmp = o.getF();
```




# Prototypes

```
function Foo(f)
{
    this.f = f;
}
```

```
Foo.prototype.getF = function()
{
    return this.f;
}
```


```
var o = new Foo(42);
```

```
var tmp = o.getF();
```




# Prototypes

```
class Foo {  
    constructor(f)  
    {  
        this.f = f;  
    }  
  
    getF()  
    {  
        return this.f;  
    }  
}  
  
var o = new Foo(42);  
  
var tmp = o.getF();
```




# Prototypes

```
class Foo {  
    constructor(f)  
    {  
        this.f = f;  
    }  
  
    getF()  
    {  
        return this.f;  
    }  
}  
  
var o = new Foo(42);  
  
var tmp = o.getF();
```



# Prototypes

```
class Foo {  
    constructor(f)  
    {  
        this.f = f;  
    }  
  
    getF()  
    {  
        return this.f;  
    }  
  
}  
  
var o = new Foo(42);  
  
var tmp = o.getF();
```




# Prototype Inline Caching

- Goals:
  - `o.getF` should be fast
  - `o.getF()` should be inlineable
  - `o.newField = value` should be fast


# Prototype Inline Caching

- Mono proto
- Transition watchpoint
- Replacement watchpoint

# Mono Proto




# Mono Proto




# Proto Chain




# Proto Chain




# Proto Chain




# Proto Chain



# Proto Chain




# Proto Chain




# Proto Chain

```
// tmp = o.getF  
if (o->structureID == 42  
    && P1->structureID == 43  
    && P2->structureID == 44)  
    return P2->inlineStorage[...]  
  
else  
    slowGet(o, "getF")
```




# Proto Chain

```
// tmp = o.getF  
if (o->structureID == 42  
    && P1->structureID == 43  
    && P2->structureID == 44)  
    return P2->inlineStorage[...]  
  
else  
    slowGet(o, "getF")
```




# Proto Chain

```
// o.newField = value
if (o->structureID == 42
    && P1->structureID == 43
    && P2->structureID == 44)
    o->inlineStorage[...] = value
    o->structureID = 100;
else
    slowPut(o, "newField",
            value)
```




# Proto Chain


```
// o.newField = value
if (o->structureID == 42
    && P1->structureID == 43
    && P2->structureID == 44)
    o->inlineStorage[...] = value
    o->structureID = 100;
else
    slowPut(o, "newField",
            value)
```




# Transition Hash Consing



# Transition Hash Consing




# Transition Hash Consing




## *Optimized Code*

```
if (o->structureID == 42  
    && P1->structureID == 43  
    && P2->structureID == 44)  
    return P2->inlineStorage[...]  
else  
    slowGet(o, "...")
```




## *Optimized Code*

```
if (o->structureID == 42  
    && P1->structureID == 43  
    && P2->structureID == 44)  
    return P2->inlineStorage[...]  
  
else  
    slowGet(o, "...")
```




## *Optimized Code*

```
if (o->structureID == 42  
    && P1->structureID == 43  
    && P2->structureID == 44)  
    return P2->inlineStorage[...]  
  
else  
    slowGet(o, "...")
```




## *Optimized Code*

```
if (o->structureID == 42  
    && P1->structureID == 43  
    && P2->structureID == 44)  
    return P2->inlineStorage[...]  
  
else  
    slowGet(o, "...")
```




### *Optimized Code*

```
if (o->structureID == 42  
    && P1->structureID == 43  
    && P2->structureID == 44)  
    return P2->inlineStorage[...]  
  
else  
    slowGet(o, "...")
```




### *Optimized Code*

```
if (o->structureID == 42)
    return P2->inlineStorage[...]
else
    slowGet(o, "...")
```




### *Optimized Code*


```
if (o->structureID == 42)
    return P2->inlineStorage[...]
else
    slowGet(o, "...")
```



## *Optimized Code*

```
if (o->structureID == 42)
    return P2->inlineStorage[...]
else
    slowGet(o, "...")
```






# Watchpoint

```
class Watchpoint {  
public:  
    virtual void fire() = 0;  
};
```


### *Optimized Code*

```
if (o->structureID == 42)
    return P2->inlineStorage[...]
else
    slowGet(o, "...")
```




### *Optimized Code*

```
if (o->structureTD == 42)
    return P2->inlineStorage[...]
else
    slowGet(o, "...")
```




### *Optimized Code*

```
if (o->structureID == 42)
    return P2->inlineStorage[...]
else
    slowGet(o, "...")
```



### *Optimized Code*

```
if (o->structureID == 42)
    return constant
else
    slowGet(o, "...")
```




# Simple Inline Caching

- `tmp = o.f // self`
- `o.f = tmp // self, existing`
- `o.f = tmp // self, new`
- `tmp = o.f // prototype`


# Simple Inline Caching

- `tmp = o.f // self`
  - `o.f = tmp // self, existing`
  - `o.f = tmp // self, new`
  - `tmp = o.f // prototype`
- 
- only need one branch*


# JIT Monomorphic IC speed-up




# JIT Monomorphic IC speed-up




# JIT Monomorphic IC speed-up




# JIT Monomorphic IC speed-up




# JIT Monomorphic IC speed-up




# JIT Monomorphic IC speed-up




# JIT Monomorphic IC speed-up




# JIT Monomorphic IC speed-up




# JIT Monomorphic IC speed-up




# JIT Monomorphic IC speed-up




# JIT Monomorphic IC speed-up



# JIT Monomorphic IC speed-up



# JIT Monomorphic IC speed-up



# Advanced Topics

- Polymorphic Inline Caches
- Inlining Inline Caches
- Type Inference

# Polymorphic Inline Caches



{f: 64, g:75}





{f: 7, g:32}



{f: 64, g:75}



{f: 98, g:23}

{f: 342, g:5}

{f: 64, g:75}

{f: 98, g:23}

{f: 7, g:32}

{f: 342, g:5}

{f: 64, g:75}

{f: 98, g:23}

{f: 7, g:32}

{f: 54, g:75,  
h:389}

{f: 6, g:18,  
h:83}


{f: 7, g:32}


{f: 64, g:75}

{f: 54, g:75,  
h:389}

{f: 98, g:23}

{f: 342, g:5}





{f: 6, g:18,  
h:83}

{g: 37, f: 30}

{f: 7, g:32}

{f, g}

{f: 64, g:75}

{f: 54, g:75,  
h:389}

{f, g, h}

{f: 98, g:23}

{g: 835, f: 0}

{g, f}


{f: 342, g:5}


```
var tmp = o.f;
```


$S1:$   
 $\{f, g\}$

$\{f: 64, g: 75\}$

`var tmp = o.f;`







```
if (o->structureID == S1 || o->structureID == S2)
    o->inlineStorage[0]
else if (o->structureID == S3)
    o->inlineStorage[1]
else
    slowGet(o, "f")
```

```
function foo(o) {  
    return o.f;  
}  
  
noInline(foo);  
  
for (var i = 0; i < 10000; ++i) {  
    foo({f: 1, g: 2});  
    foo({f: 1, g: 2, h:3});  
    foo({g: 2, f: 1});  
}
```

*S1:*  
 $\{f, g\}$

```
0x456923529be: cmp $0x129, (%rax)
0x456923529c4: jnz 0x45692352a2d
0x456923529ca: mov 0x10(%rax), %rax
0x456923529ce: nop 0x200(%rax)
```

```
function foo(o) {  
    return o.f;  
}  
  
noInline(foo);  
  
for (var i = 0; i < 10000; ++i) {  
    foo({f: 1, g: 2});  
    foo({f: 1, g: 2, h:3});  
    foo({g: 2, f: 1});  
}
```

*S1:*  
 $\{f, g\}$

*S2:*  
 $\{f, g, h\}$

*S3:*  
 $\{g, f\}$

```
0x45692352be0: mov (%rax), %esi
0x45692352be2: cmp $0x129, %esi
0x45692352be8: jnz 0x45692352bf7
0x45692352bee: mov 0x10(%rax), %rax
0x45692352bf2: jmp 0x456923529d5
0x45692352bf7: cmp $0x126, %esi
0x45692352bfd: jnz 0x45692352c0c
0x45692352c03: mov 0x10(%rax), %rax
0x45692352c07: jmp 0x456923529d5
0x45692352c0c: cmp $0x12b, %esi
0x45692352c12: jnz 0x45692352a2d
0x45692352c18: mov 0x18(%rax), %rax
0x45692352c1c: jmp 0x456923529d5
0x45692352c21: jmp 0x45692352a2d
```

*S1:*  
 $\{f, g\}$

*S2:*  
 $\{f, g, h\}$

*S3:*  
 $\{g, f\}$

```
0x45692352be0: mov (%rax), %esi
0x45692352be2: cmp $0x129, %esi
0x45692352be8: jnz 0x45692352bf7
0x45692352bee: mov 0x10(%rax), %rax
0x45692352bf2: jmp 0x456923529d5
0x45692352bf7: cmp $0x126, %esi
0x45692352bfd: jnz 0x45692352c0c
0x45692352c03: mov 0x10(%rax), %rax
0x45692352c07: jmp 0x456923529d5
0x45692352c0c: cmp $0x12b, %esi
0x45692352c12: jnz 0x45692352a2d
0x45692352c18: mov 0x18(%rax), %rax
0x45692352c1c: jmp 0x456923529d5
0x45692352c21: jmp 0x45692352a2d
```






```
0x45692352be0: mov (%rax), %esi
0x45692352be2: cmp $0x129, %esi
0x45692352be8: jnz 0x45692352bf7
0x45692352bee: mov 0x10(%rax), %rax
0x45692352bf2: jmp 0x456923529d5
0x45692352bf7: cmp $0x126, %esi
0x45692352bfd: jnz 0x45692352c0c
0x45692352c03: mov 0x10(%rax), %rax
0x45692352c07: jmp 0x456923529d5
0x45692352c0c: cmp $0x12b, %esi
0x45692352c12: jnz 0x45692352a2d
0x45692352c18: mov 0x18(%rax), %rax
0x45692352c1c: jmp 0x456923529d5
0x45692352c21: jmp 0x45692352a2d
```

The diagram illustrates a state transition process. Three states are shown as red circles: *S1* containing  $\{f, g\}$ , *S2* containing  $\{f, g, h\}$ , and *S3* containing  $\{g, f\}$ . Below the circles is a sequence of assembly instructions. A curved arrow originates from the final instruction `jmp 0x45692352a2d` and points to the word **done**.



```
0x45692352be0: mov (%rax), %esi
0x45692352be2: cmp $0x129, %esi
0x45692352be8: jnz 0x45692352bf7
0x45692352bee: mov 0x10(%rax), %rax
0x45692352bf2: jmp 0x456923529d5
0x45692352bf7: cmp $0x126, %esi
0x45692352bfd: jnz 0x45692352c0c
0x45692352c03: mov 0x10(%rax), %rax
0x45692352c07: jmp 0x456923529d5
0x45692352c0c: cmp $0x12b, %esi
0x45692352c12: jnz 0x45692352a2d
0x45692352c18: mov 0x18(%rax), %rax
0x45692352c1c: jmp 0x456923529d5
0x45692352c21: jmp 0x45692352a2d
```

***done***





```
0x45692352be0: mov (%rax), %esi
0x45692352be2: cmp $0x129, %esi
0x45692352be8: jnz 0x45692352bf7
0x45692352bee: mov 0x10(%rax), %rax
0x45692352bf2: jmp 0x456923529d5
0x45692352bf7: cmp $0x126, %esi
0x45692352bfd: jnz 0x45692352c0c
0x45692352c03: mov 0x10(%rax), %rax
0x45692352c07: jmp 0x456923529d5
0x45692352c0c: cmp $0x12b, %esi
0x45692352c12: jnz 0x45692352a2d
0x45692352c18: mov 0x18(%rax), %rax
0x45692352c1c: jmp 0x456923529d5
0x45692352c21: jmp 0x45692352a2d
```

$done$

$done$

$done$



```
0x45692352be0: mov (%rax), %esi
0x45692352be2: cmp $0x129, %esi
0x45692352be8: jnz 0x45692352bf7
0x45692352bee: mov 0x10(%rax), %rax
0x45692352bf2: jmp 0x456923529d5 → done
0x45692352bf7: cmp $0x126, %esi ←
0x45692352bfd: jnz 0x45692352c0c
0x45692352c03: mov 0x10(%rax), %rax
0x45692352c07: jmp 0x456923529d5 → done
0x45692352c0c: cmp $0x12b, %esi ←
0x45692352c12: jnz 0x45692352a2d → slow
0x45692352c18: mov 0x18(%rax), %rax
0x45692352c1c: jmp 0x456923529d5
0x45692352c21: jmp 0x45692352a2d
```




```
0x45692352be0: mov (%rax), %esi
0x45692352be2: cmp $0x129, %esi
0x45692352be8: jnz 0x45692352bf7
0x45692352bee: mov 0x10(%rax), %rax
0x45692352bf2: jmp 0x456923529d5 → done
0x45692352bf7: cmp $0x126, %esi ←
0x45692352bfd: jnz 0x45692352c0c
0x45692352c03: mov 0x10(%rax), %rax
0x45692352c07: jmp 0x456923529d5 → done
0x45692352c0c: cmp $0x12b, %esi ←
0x45692352c12: jnz 0x45692352a2d → slow
0x45692352c18: mov 0x18(%rax), %rax
0x45692352c1c: jmp 0x456923529d5 → done
0x45692352c21: jmp 0x45692352a2d
```


# Polymorphic Access JIT

- Self
- Prototype
- Transition (new property)
- Getters
- Setters
- Custom accessors
- Snippets (DOM JIT)


# JIT Polymorphic IC speed-up




# JIT Polymorphic IC speed-up




# JIT Polymorphic IC speed-up




# JIT Polymorphic IC speed-up




# JIT Polymorphic IC speed-up




# JIT Polymorphic IC speed-up




# JIT Polymorphic IC speed-up




# JIT Polymorphic IC speed-up



# JIT Polymorphic IC speed-up



# JIT Polymorphic IC speed-up




# Inlining Inline Caches

# Four Tiers




↔ *latency*                                    *throughput* →


# Four Tiers





*latency* ← → *throughput*



# Four Tiers
























$S1:$   
 $\{f, g\}$



$\text{tmp} = o.f$





















$S1:$   
 $\{f, g\}$

$tmp = o.f$   
 $tmp2 = o.g$




$S1:$   
 $\{f, g\}$

$tmp = o.f$   
 $tmp2 = o.g$




$S1:$   
 $\{f, g\}$


$tmp = o.f$   
 $tmp2 = o.g$




$S1:$   
 $\{f, g\}$

$tmp = o.f$   
 $tmp2 = o.g$







$S1:$   
 $\{f, g\}$

$tmp = o.f$   
 $tmp2 = o.g$

LLInt  
(interpreter)

```
get_by_id  
..., S1, 0
```



```
get_by_id  
..., S1, 1
```



Baseline  
(template JIT)



```
cmp S1,  
(%rax)  
jnz Lslow  
mov 10(%rax),  
%rax
```



```
cmp S1,  
(%rax)  
jnz Lslow  
mov 18(%rax),  
%rax
```



DFG  
(less optimizing JIT)


















## Inline Cache Control Flow




## Inline Cache Control Flow




## Inline Cache Control Flow




## Inlined with OSR exits



## Inline Cache Control Flow




## Inlined with OSR exits




# Benefits of Inlining

- Common subexpression elimination
- Abstract interpreter models the structure
- Object allocation sinking
- Many other optimizations


# IC Monomorphic Inlining Speed-up




# IC Monomorphic Inlining Speed-up




# IC Monomorphic Inlining Speed-up




# IC Monomorphic Inlining Speed-up




# IC Inlining and Watchpoints



# IC Inlining and Watchpoints




# IC Inlining and Watchpoints




# IC Inlining and Watchpoints




# IC Inlining and Watchpoints



# IC Inlining and Watchpoints




# IC Inlining and Watchpoints




# Minimorphic IC Inlining

```
var tmp = o.f;
```


# Minimorphic IC Inlining



# Minimorphic IC Inlining



# Minimorphic IC Inlining




CheckStructure(@o, [S1, S2])  
GetByOffset(@o, "f", 0)


# Polymorphic IC Inlining

```
var tmp = o.f;
```

# Polymorphic IC Inlining




# Polymorphic IC Inlining




**DFG IR:** MultiGetByOffset(@o, “f”, [S1, S2] => 0, [S3] => 1)


# Polymorphic IC Inlining




# IC Polymorphic Inlining Speed-up




# IC Polymorphic Inlining Speed-up




# IC Polymorphic Inlining Speed-up




# IC Polymorphic Inlining Speed-up




# IC Polymorphic Inlining Speed-up



# IC Polymorphic Inlining Speed-up



# IC Polymorphic Inlining Speed-up





```
function foo(o) { return o.f; }
```


$S_1$ :  
 $\{f, g\}$

$S_2$ :  
 $\{f, g, h\}$


$S_3$ :  
 $\{g, f\}$

```
function foo(o) { return o.f; }
```


```
function bar(p) { return foo(p.g); }
```



```
function bar(p) { return foo(p.g); }
```




```
function bar(p) { return foo(p.g); }
```




```
function foo(o) { return o.f; }
```



```
function bar(p) { return foo(p->g); }
```



```
function bar(p) { return foo(p->g); }
```



DFG  
(less optimizing JIT)

...

→ foo

```
cmp S1,  
     (%rax)  
jnz Lslow  
mov 10(%rax),  
     %rax
```

<- foo



DFG  
(less optimizing JIT)



...

→ foo


```
cmp S1,
     (%rax)
jnz Lslow
mov 10(%rax),
     %rax
```

<- foo


FTL  
(optimizing JIT)




# IC Polyvariant Inlining Speed-up




# IC Polyvariant Inlining Speed-up



# IC Polyvariant Inlining Speed-up




# IC Polyvariant Inlining Speed-up




**DFG IR:** MultiGetByOffset(@o, “f”, [S1, S2] => 0, [S3] => 1)  
MultiGetByOffset(@o, “g”, [S1, S2] => 1, [S3] => 0)


**DFG IR:**


```
MultiGetByOffset(@o, "f", [S1, S2] => 0, [S3] => 1)
MultiGetByOffset(@o, "g", [S1, S2] => 1, [S3] => 0)
```





**B3 IR:**








































B3 IR before:




B3 IR before:




**duplicateTails(procedure);  
foldPathConstants(procedure);**


B3 IR after:




# Taildup Speed-up




# Taildup Speed-up



# Taildup Speed-up




# Taildup Speed-up



$\circ \cdot f \cdot g$

*o.f.g*



- a: CheckStructure(@o, S1)
- b: GetByOffset(@o)
- c: CheckStructure(@b, S2)
- d: GetByOffset(@c)

`o.f.g`



`a: CheckStructure(@o, S1)`

`b: GetByOffset(@o)`

`c: CheckStructure(@b, S2)`


`d: GetByOffset(@c)`

# Property Type Inference


# Property Type Inference

{f: , ...}


# Property Type Inference




# Property Type Inference




# Property Type Inference




# Property Type Inference




# Property Type Inference




# Property Type Inference



# Property Type Inference




# Property Type Inference




- Stores to f check type
- Loads don't have to
- S1's inferred type table  
watches S2's transitions


# Inferred Type “Speed-up”




# Inferred Type “Speed-up”




# Inferred Type “Speed-up”




# Inferred Type “Speed-up”




# Inferred Type “Speed-up”



# Inferred Type “Speed-up”




```
function foo()
{
    class Helper {
        ...
    }
    var h = new Helper();
    ...
}
```




```
function foo()
{
    class Helper {
        ...
    }
    var h = new Helper();
    ...
}
```


```
function foo()
{
    class Helper {
        ...
    }
    var h = new Helper();
    ...
}
```




```
function foo()
{
    class Helper {
        ...
    }
    var h = new Helper();
    ...
}
```




```
function foo()
{
    class Helper {
        ...
    }
    var h = new Helper();
    ...
}
```




```
function foo()
{
    class Helper {
        ...
    }
    var h = new Helper();
    ...
}
```




```
function foo()
{
    class Helper {
        ...
    }
    var h = new Helper();
    ...
}
```




## Mono Proto




## Poly Proto object structure




# Poly Proto Speed-up




# Poly Proto Speed-up



# Poly Proto Speed-up



# Poly Proto Speed-up



# How I Ran These Experiments

- WebKit trunk revision 233406
- Safari trunk hash  
`dd8296dbaac7afa5ed9a699aa261033ea5f5577c`
- macOS Internal SDK
- make release (not a root, no shared cache)
- Patch, scripts, and raw data at [webkit.org/b/187414](https://webkit.org/b/187414)
- MacBookPro11,5 2.8GHz 16GB RAM 1TB SSD

